

Answer 1


✓ CORRECT ANSWER



A WINS: S3 — Object Storage

The right choice for this file upload service.

System Design Labs



Massive Scale, No Limits

Handles 10TB today, 100TB+ tomorrow without any changes.



Cost Efficient at Scale

At 100TB → ~\$2,300/month. Pay only for what you use.



Lifecycle to Glacier

Automatically move cold files to Glacier. Save up to 90%.



Event Driven

S3 events trigger your ML pipeline the moment a file lands.



Built-in Security & Compliance

Versioning, encryption, audit logs — all built in, zero extra work.



For greenfield cloud services, this is the default.

Only deviate when you have a specific reason not to.

A

S3 — object storage, infinite scale, pay per GB, no servers to manage.

B

EBS — block storage, SSD-backed, attach to your EC2, low latency reads.

C

EFS — managed NFS, shared across multiple EC2 instances simultaneously.

D

MinIO on EC2 — S3-compatible self-hosted object storage, you own the infra.

Why A wins (S3):

One storage layer. Zero capacity planning. Every service reads from the same bucket using the S3 key stored in your DB.

At 100TB → ~\$2,300/month. Lifecycle policies move cold files to Glacier automatically. S3 events trigger your ML pipeline the moment a file lands. Versioning, encryption, audit logs — all built in, zero extra work.

For greenfield cloud services, this is the default. You only deviate when you have a specific reason not to.

Answer 2

Why B destroys your architecture quietly (EBS):

EBS is a disk. A fast, reliable disk — attached to **one EC2 instance at a time**.

The moment your ML pipeline runs on a different instance than your upload service, it can't read those files. Add auto-scaling? Each new instance gets a blank disk. Deploy across two AZs? Same problem.

Works perfectly in dev. Silently breaks the moment you scale horizontally. The team that picks EBS here usually discovers this 6 months later when they're doing a painful migration under pressure.

Answer 3

Why C is great — but not here (EFS):

EFS is genuinely powerful. Shared NFS, multiple EC2 instances mount it simultaneously, full POSIX semantics.

The problem: **\$0.30/GB/month**. At 100TB that's **\$30,000/month** vs S3's **\$2,300/month**. A 13x cost difference.

EFS is the right answer when your legacy app is hardcoded to `open()` / `read()` / `write()` filesystem calls and can't be refactored. For a greenfield NestJS service with an S3 SDK? You're paying 13x for filesystem semantics you don't need.

Answer 4


Why D is a cost-savings trap (MinIO):

MinIO is a real product used at hyperscale. S3-compatible API, high throughput, runs on your own hardware.

But self-hosted means **you own availability, disk failures, backups, capacity planning, and 3am incidents**. At 10–100TB, the engineering cost of running MinIO reliably exceeds S3 spend by a wide margin. You're not saving money — you're converting cloud spend into engineering hours, which cost more.

MinIO makes sense at petabyte scale with dedicated infra teams and painful egress bills. Not here.

Cutting Cross-Region Latency



Your e-commerce platform just crossed 2M daily active users. 70% are in the US, 30% in Europe.

Latency complaints are piling up from European users — 380ms average round-trip to your US-East region. Support tickets are up 40%. Black Friday is in 6 weeks.

Your infrastructure: single AWS us-east-1 region, RDS PostgreSQL (primary), Redis cache, 12 microservices behind an API Gateway.

You need to get European latency under 80ms. The engineering team is debating four approaches.

Here's your constraint: you cannot afford a full database rewrite, and you need this shipped before Black Friday.

A) Active-Active multi-region — deploy the full stack in eu-west-1, use a distributed database (CockroachDB or Aurora Global), route users to the nearest region. Writes go to both regions simultaneously.

B) Active-Passive with read replicas — keep us-east-1 as primary, spin up eu-west-1 as a hot standby with read replicas. European reads go local, writes still go to US. Failover in minutes if US goes down.

C) CDN + Edge caching — keep the single region, push static assets and cacheable API responses to CloudFront edge nodes in Europe. No database changes.

D) Active-Active with eventual consistency — deploy full stack in both regions, allow each region to own its writes, sync asynchronously. Accept that a European user might see a US write 200ms late.


Answer 1

✓ CORRECT ANSWER

Why B wins: Active-Passive with read replicas

System Design Labs

Fastest way to get European read latency under 80ms — without rewriting the database.



Why B wins:

The constraint is the answer. You can't rewrite the database, and you have 6 weeks.

European users are complaining about *read* latency — browsing products, checking order status, loading their profile. Those are reads. An RDS read replica in eu-west-1 costs you 1-2 days of infra work and immediately moves European reads to ~15ms instead of 380ms. Writes still go to us-east-1 (acceptable — users tolerate slightly higher write latency for checkout). Replication lag is typically under 100ms.

Active-passive also gives you a real failover story: US goes down → promote EU replica in minutes.

Answer 2

Why A is the trap (Active-Active + distributed DB):

This is the "correct at scale" answer that kills teams before Black Friday. Migrating from RDS PostgreSQL to CockroachDB or Aurora Global under a 6-week deadline is a multi-quarter project — distributed transactions, global clock skew, new failure modes to test. Right architecture for the long term. Wrong answer for the problem as stated.

Answer 3

Why C is wrong (CDN + Edge caching):


CDN solves static asset latency — images, JS, CSS. It does nothing for dynamic reads (user cart, order history, account data) and zero for write latency. It's a layer you add *on top of* a real solution, not instead of one. Partial credit, not the answer.

Answer 4

Why D is dangerous (Active-Active + eventual consistency):

Most honest version of multi-region — and the most dangerous to deploy without prep. Eventual consistency across regions means a European user places an order and your US region sees a different inventory count for 200ms. For limited stock = oversell. For checkout = duplicate charges. Conflict resolution is a product problem disguised as an infra problem.

Managing Secrets and Credentials



Your startup just got its first SOC 2 audit.

The auditor asks: "Where are your database passwords, API keys, and service tokens stored?"

Your senior engineer goes quiet.

Turns out half of them are in .env files committed to git 18 months ago. Three are hardcoded in Lambda environment variables. One is in a Slack message from 2023.

You have 6 services in production, 4 environments, and zero rotation policy.

Here's the setup:

- NestJS API → Postgres (password in env var)
- NestJS API → Stripe (API key in env var)
- Background workers → SQS, S3 (AWS credentials in env var)
- 3rd-party webhooks → HMAC secrets in env var
- Zero rotation. Zero audit trail. Zero centralized access control.

You need to fix this. And you can't take downtime.

- A)** Move everything to AWS Secrets Manager — SDK calls at runtime, IAM controls access, auto-rotation built in.
- B)** Use HashiCorp Vault — dynamic secrets, fine-grained policies, works across any cloud or on-prem.
- C)** Use environment variables injected at deploy time via CI/CD — secrets stored in GitHub Actions / GitLab CI secrets vault, never touch disk.
- D)** Encrypt secrets with KMS and store ciphertext in your own database — decrypt at runtime, full control.

Answer 1

✓ CORRECT ANSWER



A — AWS Secrets Manager (correct)

You're already AWS-native (Lambda, SQS, S3). Secrets Manager gives you:

- Centralized encrypted storage (KMS-backed)
- IAM-based access control per secret per service
- Native auto-rotation for RDS + custom Lambda rotators for everything else
- Every access logged in CloudTrail → exactly what your SOC 2 auditor wants

Migration is zero-downtime: swap env vars for `getSecretValue()` calls, deploy, rotate the old creds after confirming the new ones work. Cost: ~\$0.40/secret/month. For 20 secrets that's \$8/month.

Answer 2

B — HashiCorp Vault (senior engineer trap)

Vault is *more powerful* — dynamic secrets (a fresh Postgres credential per-request that expires in 1h) is one of the most elegant patterns in production infra.

But Vault is infrastructure you own and operate. You need a Raft cluster (3+ nodes), a storage backend, unsealing procedures, DR planning, and an ops team that knows Vault. If Vault goes down, secrets become inaccessible — services fail.

When Vault wins: Multi-cloud, on-prem requirements, dynamic secrets at scale, or you have a dedicated infra team. If you're AWS-native under 50 engineers — Secrets Manager first, always.

Answer 3

C — CI/CD Injected Env Vars (partially right, wrong scope)

GitHub Actions / GitLab CI secrets are the right place for *build-time* and *deploy-time* secrets — Docker Hub creds, Terraform tokens, deploy keys.

They're wrong for *runtime* secrets — the DB password your app needs at 3 AM, the Stripe key for a live webhook. Those get injected as env vars into ECS task definitions or Lambda config — harder to rotate atomically, no audit trail on runtime access.

C complements A or B. It doesn't replace them.


Answer 4

D — DIY KMS + own database (wrong)

This is reinventing Secrets Manager badly. You now own: encryption key management, rotation logic, access control, audit logging, and a database that itself needs credentials to connect to.

The meta-problem: what secret protects the secret store? You end up with a bootstrapping problem that managed secret stores are specifically designed to solve. This is what teams built in 2012 before managed options existed.

Reconstructing State with Event Sourcing



Your order service takes 200 writes/sec at peak.

You audit 6 months of data. Something's off — two orders show the same ID, different totals.

You have the current state. You don't have how it got there.

Your DB is a graveyard of overwritten rows.

Here's the system:

- OrderService → Postgres (current state only)
- Events: placed, updated, cancelled, refunded
- Every UPDATE overwrites the previous row
- No audit log. No event history. No replay.

A billing dispute just landed. You need to reconstruct exactly what happened to Order #8471. You can't.

Instead of storing the current state, you store the sequence of events that produced it.

What's your approach when redesigning this service?

- A** Event Sourcing — append-only event log as the source of truth, current state derived from replaying events.
- B** Change Data Capture (CDC) — keep Postgres as-is, but stream all row changes to Kafka for an audit trail.
- C** Add an audit_log table — trigger-based shadow writes on every INSERT/UPDATE/DELETE.
- D** Dual-write — write to both the current-state table and a separate events table on every operation.


Answer 1

✓ CORRECT ANSWER

ANSWER: A — EVENT SOURCING ✓

System Design Labs

Append-only event log as the source of truth; current state derived from replaying events.

**A EVENT SOURCING**

Append-only event log as the source of truth; current state derived from replaying events.

- ✓ Full replay & time travel
- ✓ Projection flexibility
- ✓ Immutable & auditable
- ✓ True source of truth

B CHANGE DATA CAPTURE (CDC)

Keep Postgres as-is, but stream all row changes to Kafka for an audit trail.

- ✗ Misses in-app context and intent
- ✗ Hard to reconstruct business events
- ✗ Lag, ordering & schema evolution challenges
- ✗ It's still a side channel, not the source of truth

C ADD AN AUDIT_LOG TABLE

Trigger-based shadow writes on every INSERT/UPDATE/DELETE.

- ✗ Easy to bypass or break
- ✗ Adds write overhead to hot path
- ✗ Captures data changes, not business events
- ✗ Schema coupling & maintenance burden

D DUAL-WRITE

Write to both the current-state table and a separate events table on every operation.

- ✗ Dual-write inconsistency risk
- ✗ Requires distributed transactions or outbox
- ✗ Complex failure modes
- ✗ Not guaranteed sequence or completeness

Answer: A — Event Sourcing ✓

Here's why, and why the other three look reasonable but miss the point:

Why A wins (Event Sourcing):

Event Sourcing flips the model entirely. Instead of storing the *latest state* and losing how you got there, you store every event that ever happened. The current state is a projection — computed by replaying events from the beginning (or from a snapshot).

For Order #8471: you don't query a row. You replay OrderPlaced, OrderUpdated, PaymentCaptured, RefundInitiated. Every mutation is preserved, timestamped, immutable. You can reconstruct state at any point in time. You can build new projections (e.g., "total revenue by SKU last 90 days") without changing the core model — just replay and project differently.

This is what Axon, EventStoreDB, and the event-sourced layers in most serious fintech/e-commerce systems do. It's not a log — it's the primary data model.

Answer 2

Why B is the trap (CDC):

CDC is powerful and often underrated. Tools like Debezium + Kafka let you stream every Postgres row change downstream without touching your app. It looks like Event Sourcing from the outside — you have a sequence of changes.

But here's the trap: CDC captures *state deltas*, not *business events*. You get total changed from 89.99 to 79.99 — not DiscountApplied by coupon SAVE10. The semantic meaning is lost. You can reconstruct what changed, not *why* it changed. That distinction kills you during disputes and compliance audits.

CDC is a great complement to Event Sourcing for propagation. It's not a substitute for it.

Answer 3

Why C is wrong (audit_log table):

Trigger-based audit logs are the classic band-aid. The problem: you're treating the row as source of truth and logging changes as a side effect.

Side effects break. Triggers get disabled during migrations. Bulk imports skip them. Schema changes orphan the trigger. Six months later, the audit_log has gaps — exactly when you need it most. It's a shadow, not a foundation.


Answer 4

Why D is wrong (Dual-write):

Dual-write seems reasonable. But you've created a distributed consistency problem inside your own service. If the state write succeeds and the events write fails — network hiccup, crash, OOM — you've got state with no corresponding event. Audit trail unreliable by design.

This is the exact problem the Outbox Pattern was invented to solve. But if you're going that far, you're halfway to proper Event Sourcing anyway.

Improving LLM Classification Accuracy



Your AI feature works in the demo.

It fails in production 3 weeks later. Nobody touched the model. Nobody changed the code.

The only thing that changed: the inputs got messier.

Here's the setup:

You're at a SaaS company. 50,000 support tickets a week. Your team builds an AI triage system — GPT-4o classifies each ticket into 6 categories (billing, bug, feature request, account access, security, other) so the right team gets it instantly.

In dev, it nails 71% accuracy. You need 90%+ to cut manual review.

The model is locked. The budget for inference isn't unlimited. You need to close the 19-point gap.

Here are your four options:

A) Zero-shot with a better system prompt — rewrite the instructions, add explicit category definitions, specify edge case rules. No examples.

B) Few-shot examples — add 3–5 real classified tickets directly in the prompt. One example per category edge case.

C) Chain-of-Thought — add "think step-by-step before answering" to the prompt. Force the model to reason through the ticket before outputting the category.

D) Self-Consistency — run each ticket through the model 5 times with temperature=0.7, take the majority vote across outputs.

Answer 1

✓ CORRECT ANSWER

B — Few-shot Best answer here

Accuracy: 88–93%. Cost increase: near-zero.

Few-shot shifts the model from *reading rules* to *pattern-matching against real examples*. Fundamentally stronger signal.

The part most teams get wrong: they pick the obvious examples. A crystal-clear billing ticket. A textbook bug report. Those teach the model nothing it didn't already know.

You want the *ambiguous* ones. The ticket that looks like billing but is account access. The feature request that reads like a bug. Those are where the model currently fails — and those are the examples that close the gap.

Get the selection right, you close most of the 19 points without touching anything else.

Answer 2**A — Zero-shot with a better system prompt**

Accuracy ceiling: ~78–82%.

This is where every team starts. You iterate the prompt, add category definitions, enumerate edge cases. It feels like progress. And for a while it is.

The hard wall: instructions *describe* categories. They can't *show* the model where the ambiguous cases land.

Real example: you define "billing" as anything about charges or invoices. Then you get this ticket — "*I can't log in and I was charged twice last month.*" Billing? Account access? Both?

No amount of rewriting tells the model which bucket that gets. It guesses wrong 20–25% of the time on edge cases. And your dataset is full of edge cases. The ceiling is real.

Answer 3

C — Chain-of-Thought

Accuracy delta: $\pm 2\%$. Latency: +200–400ms.

CoT is the most misapplied technique in production AI right now. Teams add "think step-by-step" to everything. For some tasks it works. Not this one.

Why it struggles here: the model generates a reasoning chain, then commits to a label. On ambiguous tickets, that chain becomes a liability. "*Mentions a charge... but also a login issue... charge is mentioned first... probably billing.*" It picks billing. Right answer is account access.

CoT helps when the answer *emerges from reasoning* — math, debugging, planning. For classification, the answer emerges from calibration. Reasoning doesn't substitute for calibration.

Layer it on top of few-shot if you're still missing edge cases. Don't reach for it first.

Answer 4

D — Self-Consistency

Accuracy: 93–96%. Cost: 5x. Latency: 5x.


Run the same ticket 5–10x at temperature=0.7, take the majority vote. Works because each run samples a slightly different path — noise cancels out across runs.

The math is good. The cost is brutal.

50K tickets/week $\times$ \$0.002 $\times$ 5 runs = **\$500/week** vs **\$100/week** with few-shot. You're paying \$400/week extra for a 3–5% gain on the hardest cases.

Worth it when being wrong is expensive: fraud detection (a missed flag costs hundreds), medical triage (wrong category = real harm). For support routing, a misclassified ticket costs ~\$0.10 of human rerouting time. Use B.

Geospatial “Find Nearby Drivers” at Scale



You're building the "find nearby drivers" feature for a ride-hailing app.

At peak, you have 500,000 active drivers updating their GPS location every 5 seconds. Riders query for drivers within 2km. At scale, you're doing $\sim 100,000$ proximity queries per second.

Your naive implementation does this:

```
SELECT * FROM drivers
```

```
WHERE lat BETWEEN ? AND ?
```

```
AND lng BETWEEN ? AND ?
```

It works fine at 1,000 drivers. At 500,000, it's a full table scan on every query. Latency hits 800ms. Riders see a spinner. Drivers miss trips.

Your team proposes 4 approaches to fix this:

A) Geohash partitioning — Encode each driver's location into a geohash string. Index by geohash prefix. Proximity queries become a string lookup on the index.

B) PostGIS with spatial indexes — Add a PostGIS extension to Postgres. Use a proper R-tree/GiST spatial index for bounding-box and radius queries.

C) Quadtree in memory — Keep all active driver positions in a quadtree data structure in a Redis-backed in-memory service. Decompose space recursively until each cell has $\leq N$ drivers.

D) H3 hexagonal grid (Uber's system) — Divide the earth into hexagonal cells at multiple resolutions. Assign each driver to a cell. Queries check the target cell + 6 neighbors at the right resolution.

You need sub-50ms p99 latency, real-time updates, and it has to stay accurate at cell boundaries.


Every option here ships with its own diagram, so no single image marks the answer; the author's pick is **D — H3 Hexagonal Grid** (Answer 1 below).

Answer 1 ✓ CORRECT ANSWER

ANSWER: D — H3 Hexagonal Grid

System Design Labs

Best for accurate, low-latency "find nearby drivers" at scale.



★ **BOTTOM LINE:** H3 gives accurate, boundary-safe, multi-resolution proximity search with sub-10ms latency at massive scale.

Answer: D — H3 Hexagonal Grid

Why D wins:

Uber open-sourced H3 for exactly this problem. Hexagons have 6 neighbors, all equidistant from the center. Squares (geohash) have corner neighbors that are ~40% farther — this distorts "within 2km" queries. H3 also has 16 resolution levels (no reindexing on zoom changes). The killer feature: the k-ring query (target cell + 6 neighbors) guarantees no driver gets missed near a boundary. Uber does sub-10ms lookups at millions of drivers: get cells in ring → fetch IDs from Redis set keyed by cell. Lyft, Airbnb, DoorDash all use variants.

Answer 2

ANSWER: A is the trap — Geohash

System Design Labs

Production-grade, widely used — and wrong for real-time, sub-km proximity.



Boundary misses

Drivers 1m apart can land in completely different prefixes.



Z-order curve problem

Geohash uses interleaved bits (Z-order). Adjacent prefix strings ≠ adjacent geography.



Systematic miss rate

At scale (100K QPS), boundary misses become a real, measurable problem.



Good for partitioning, not proximity


Great for coarse sharding by region. Wrong for real-time sub-km lookup.



Latency suffers

More cells to cover + post-filtering = higher latency and cost.

HOW GEOHASH LOOKUP WORKS (AND FAILS)



BOUNDARY MISS EXAMPLE



Same prefix lookup won't return both.

One driver gets missed.
This happens everywhere at cell boundaries.

WHY A LOSES TO OTHER APPROACHES



Not boundary-safe

Z-order (geohash) doesn't align with circular distance. Edges and corners cause misses.



Unpredictable coverage

To cover 2km, you need a variable number of prefixes (4 to 40+). More cells = more reads + filtering.



Higher latency at scale

More prefixes → more data read → more post-filtering. Hard to hit sub-50ms p99.



No built-in neighbor guarantee

You must manually compute which adjacent prefixes to include. Easy to get wrong, still not perfect.



Wrong tool for the job

Great for coarse partitioning / sharding / caching by region. Wrong for real-time, accurate, sub-km proximity search.



Bottom line: Geohash is the trap. It looks simple, works at low scale, but guarantees misses at the boundaries. At 100K QPS with sub-2km accuracy requirements, it will hurt riders and drivers.



Use a boundary-safe approach like H3 (hex grid) instead.

Why A is the trap (Geohash):

Production-grade, widely used — and wrong here. Geohash uses a Z-order curve (interleaved bits) that doesn't guarantee geographic adjacency for adjacent prefix strings. Two drivers 1 meter apart can have entirely different prefix strings if they straddle a cell boundary. This becomes a systematic miss rate at 100k queries/sec. Geohash is right for coarse partitioning (sharding by region), wrong for real-time sub-km proximity.

Answer 3

ANSWER: B fails at scale — PostGIS

System Design Labs

PostGIS is correct — just not for 500k drivers updating every 5 seconds.



Write-heavy workload
100,000 writes/sec
(500k drivers × 12 updates/min)



GiST index is write-expensive
Every write updates multiple index pages.



Database burns capacity on index maintenance
Less headroom for queries and growth.




Hard to scale writes
Postgres/PostGIS doesn't scale writes horizontally.



Latency suffers under load
Contention, VACUUM, bloat, and lock amplification.

THE FLOW (AND WHERE IT BREAKS)



WHY THIS DOESN'T SCALE

- ❌ **GiST index write amplification**
GiST isn't optimized for high write rates. Each insert/update touches multiple index pages.
- ❌ **Database becomes write-bound**
At 100k writes/sec, most CPU, IOPS, and memory go to index maintenance, not serving queries.
- ❌ **Vertical scaling hits a wall**
Postgres/PostGIS scales up, not out. Bigger box = temporary relief, not a solution.
- ❌ **Operational overhead grows**
VACUUM, REINDEX, bloat management, autovac tuning — constant firefighting.
- ❌ **Not built for real-time, global, massive scale**
Great for thousands to tens of thousands. Not for 500k drivers updating every 5 seconds with 100k QPS queries.

GOOD FOR THIS (NOT THIS)

- ✅ **Good fit**
 - Fleet dashboard
 - 100s to 5,000 assets
 - Ad-hoc spatial analytics
 - Internal tools
- ❌ **Not a fit**
 - 500k active drivers
 - 100k writes/sec
 - 100k QPS proximity queries
 - Sub-50ms p99 requirement

POSTGIS IS THE STARTING POINT, NOT THE DESTINATION



- ⚠️ BOTTOM LINE**
- PostGIS is the right tool for correctness and small/medium scale. But at 500k drivers with 100k writes/sec, it becomes the bottleneck. You'll spend more time fighting the database than building features.
- Use PostGIS to validate. Use H3 + Redis to scale.**

Why B fails at scale (PostGIS):

PostGIS is correct — it's the right answer for 5,000 trucks on a fleet dashboard. Not for 500k drivers updating every 5s. The math: $500k \times 12 \text{ updates/min} = 100k \text{ writes/sec}$ against a GiST index. GiST indexes are write-expensive. Your DB burns half its capacity on index maintenance. And you can't easily scale Postgres writes horizontally. PostGIS is the starting point — H3 + Redis is where you land when it breaks.

Answer 4

ANSWER: C almost works — Quadtree

System Design Labs

Great in theory ($O(\log n)$). Breaks under ride-hailing scale and real-world density.

 **$O(\log n)$ operations**
Insert, delete, and query are fast in a balanced tree.

 **Dynamic & adaptive**
Subdivides space only where drivers exist. Saves memory.


 **High churn hurts**
Drivers cross cell boundaries constantly → delete + reinsert + potential rebalancing.

 **Same boundary problem**
Rectangular cells create edge and corner issues. Drivers near borders get missed.


 **Uneven in dense cities**
Heavy hotspots create deep, unbalanced trees → more hops, more latency.

 **Not a complete solution**
A quadtree is a building block, not production-ready for 100k QPS and 500k moving drivers.

HOW A QUADTREE WORKS (AND WHERE IT HURTS)



BOUNDARY PROBLEM (RECTANGULAR CELLS)



 **GOOD FOR**

- Low to medium scale
- Static or low-churn data
- Games, simulations, robotics
- Spatial databases (read-heavy)

 **NOT GOOD FOR**

- 500k moving drivers
- 100k writes/sec
- Strict sub-50ms p99 SLA
- Real-time proximity at scale

THE MATH

500,000 drivers × 12 updates/min
= **100,000 writes/sec**

Delete + Insert + (Maybe) Rebalance
→ Millions of node ops/sec
→ CPU + memory + GC pressure

WHY C ALMOST WORKS (BUT DOESN'T AT THIS SCALE)

-  **High write amplification**
100,000 writes/sec → 100,000 deletes + inserts/sec on average. Rebalancing adds more CPU and memory churn.
-  **Boundary churn**
Drivers jitter around cell edges. Small GPS noise → unnecessary move operations → extra writes.
-  **Rectangles distort distance**
Square cells have farther corners (~41% farther) than centers. Harder to guarantee "within 2km" accuracy.
-  **Hotspot imbalance**
Dense downtown areas produce deep trees. Query becomes $O(\log n + k)$ with large k .
-  **Operational complexity**
Custom in-memory structure, shard management, failover, rebuilds, snapshots — a lot to get right.
-  **Not horizontally scalable by nature**
You still need smart sharding + data movement. Complex to scale writes linearly.

 **BOTTOM LINE**


Quadtree has the right idea: adaptive spatial partitioning. But at ride-hailing scale, write churn, boundary issues, and imbalance make it too expensive and complex.

It's a building block — not the full solution. H3 + Redis wins.

Why C almost works (Quadtree):

$O(\log n)$ insert/query, great in theory. Breaks at ride-hailing density: driver crossing a cell boundary = delete + reinsert + potential rebalancing, 100k times per second. Same rectangular boundary problem as geohash. Dense cities create deep uneven trees. A quadtree is a building block — not a complete solution at this scale.

Choosing HTTP/3 vs HTTP/2 at the Edge



Your API response went from 320ms to 95ms after a CDN switch.

Nothing changed on the server. Same origin. Same payload.

The only difference: the CDN started speaking HTTP/3 to your clients.

Here's the setup:

Mobile app → Load Balancer → Origin (NestJS) → DB

Every request goes through the CDN first. Latency is acceptable — until high packet-loss environments (mobile 4G → 5G transitions, flaky WiFi, Asia-Pacific routes). You're seeing 800ms+ tail latency for the p99. The CDN supports HTTP/2 and HTTP/3. What do you enable?

- A** HTTP/2 only — multiplexing over a single TCP connection eliminates the old HTTP/1.1 head-of-line problem. Proven, widely supported.
- B** HTTP/3 only — built on QUIC (UDP), eliminates TCP head-of-line blocking entirely, 0-RTT connection resumption. Modern clients handle it fine.

C HTTP/2 to origin, HTTP/2 to clients — maximize multiplexing end-to-end, avoid QUIC instability in enterprise firewalls.


D HTTP/3 to clients, HTTP/2 to origin — QUIC handles the lossy last mile, HTTP/2 handles the stable datacenter leg.

Answer 1

✓ CORRECT ANSWER

ANSWER: D — HTTP/3 to clients, HTTP/2 to origin ✓

System Design Labs



CLIENT SIDE (LAST MILE) — HTTP/3 over QUIC	ORIGIN SIDE (DATACENTER) — HTTP/2 over TCP	WHY D IS CORRECT
Packet loss: 2% HTTP/2 over TCP Stream 1, Stream 2, ..., Stream 20. One lost packet stalls ALL streams (Transport-layer HOL blocking). VS HTTP/3 over QUIC Stream 1, Stream 2, ..., Stream 20. One lost packet stalls ONLY that stream (Other streams keep flowing). 0-RTT Resumption Saves 1–2 round trips on every reconnect	Stable, low-loss private network QUIC's reliability advantages disappear. Mature, well-optimized HTTP/2 is proven and highly efficient. Lower CPU overhead QUIC encrypts everything at the transport layer — no cleartext, ever. Datacenter firewalls and tooling Better compatibility with existing enterprise infrastructure.	The problem is the last mile. TCP's transport-layer HOL blocking kills performance on lossy mobile connections. HTTP/3 (QUIC) to clients eliminates transport-layer HOL blocking, keeps other streams moving. HTTP/2 to origin leverages a stable network and avoids QUIC's crypto CPU overhead. This is the production pattern used by Cloudflare, Fastly, and Akamai.

Industry Standard: Major CDNs use HTTP/3 on the edge-facing side and HTTP/2 on the origin-facing side. It delivers the best performance where it matters, with the right tradeoffs where it doesn't.

Answer: D — HTTP/3 to clients, HTTP/2 to origin

The problem is the last mile. Mobile clients on lossy connections are where TCP falls apart.

HTTP/2 fixed HTTP/1.1's application-layer head-of-line blocking by multiplexing streams over one TCP connection. But TCP itself still has head-of-line blocking at the transport layer. A single lost packet stalls ALL streams until it's retransmitted. On a 4G connection dropping 2% of packets, your 20-stream HTTP/2 connection grinds to a halt.

HTTP/3 runs on QUIC (UDP). Each stream is independent. A dropped packet stalls only that stream, the other 19 keep flowing. Add 0-RTT resumption (subsequent connections skip the handshake entirely) and you shave 1-2 round trips on every reconnect. That's where the latency gains come from.

But your datacenter leg (CDN to origin) is a different story. It's a stable, low-loss private network. QUIC's benefits disappear on reliable connections. HTTP/2 multiplexing is mature, well-optimized, and doesn't carry QUIC's CPU overhead for cryptography (QUIC encrypts everything at the transport layer, no cleartext, ever).

This is exactly how Cloudflare, Fastly, and Akamai operate: HTTP/3 on the edge-facing side, HTTP/2 (or even HTTP/1.1 with keepalive) on the origin-facing side.

Answer 2

Why A falls short (HTTP/2 only):

HTTP/2 multiplexing is a genuine improvement over HTTP/1.1, but it doesn't solve TCP head-of-line blocking. On lossy mobile connections, you're still at the mercy of the TCP retransmission window. For p99 tail latency on bad networks, HTTP/2 won't move the needle.

Answer 3

Why B is wrong (HTTP/3 only):


HTTP/3 to origin introduces real operational risk. Enterprise firewalls and some cloud load balancers block UDP port 443 or rate-limit it. Your origin infra isn't the bottleneck, the last mile is. Fixing the origin leg with QUIC doesn't help the problem; it just adds complexity and potential breakage.

Answer 4

Why C is wrong (HTTP/2 end-to-end):

Same problem as A. You've done nothing about the last mile. Consistent HTTP/2 end-to-end sounds clean, but "clean" doesn't show up in your p99.

Multi-Writer Conflicts in Collaborative Editing



Two users edit the same document at 9:03 AM.

No locking. No coordination. Just two clients writing to the same record.

Both changes hit your server 300ms apart. One lands first. The other overwrites it.

User A's work is gone. They never get an error. They just... lose their change.

This is the multi-writer conflict problem. Every collaborative system hits it eventually. Your choices:

A) Last-Write-Wins (LWW) — highest timestamp takes the record. Loser's change disappears silently.

B) Vector Clocks — track causality per replica, detect conflicting versions, surface them to the application for resolution.

C) CRDTs (Conflict-free Replicated Data Types) — data structures that mathematically merge concurrent writes without coordination.

D) Operational Transformation (OT) — transform each operation relative to concurrent operations before applying, keeping intent intact.

Answer 1


✓ CORRECT ANSWER

C

CRDTs wins at scale

System Design Labs

Any two replicas can always be merged — mathematically, without coordination.



Answer: C — CRDTs (but the real answer depends on your conflict model — full breakdown below)

C — CRDTs wins at scale

CRDTs are data structures designed so any two replicas can always be merged — mathematically, without coordination. The key insight: if operations are commutative, associative, and idempotent, conflicts structurally *can't* exist.

Examples:

- **G-Counter:** only increments. Merge = take max per replica. Always converges.
- **LWW-Element-Set:** union of adds + removes with timestamps.
- **RGA / YATA (sequence CRDTs):** each character gets a unique ID, concurrent inserts are deterministically ordered — this is collaborative text editing.

Figma's canvas runs on CRDTs. Notion's block model is CRDT-based. At millions of concurrent writers across distributed regions, every node converges independently — zero coordination required.

The catch: CRDTs constrain your data model. Not everything maps cleanly. Some operations (like moving a file between directories) are notoriously hard to express without losing intent.

Answer 2

D — Operational Transformation (OT) ⚠️ impressive but hard

OT was Google Docs' original approach. When two operations arrive concurrently, you transform each one relative to the other before applying it. User A inserts at position 5, User B deletes at position 3 simultaneously — OT adjusts A's insert position to account for B's delete. Intent preserved.

That's why Google Docs felt magical in 2009.

But OT is operationally brutal. The transformation functions are notoriously hard to get right — the original OT papers had bugs. Every new operation type adds transformation pairs that grow quadratically. OT also requires a centralized server to sequence operations; it doesn't distribute cleanly.

Modern Google Docs has moved toward a hybrid incorporating CRDT-like ideas. OT alone doesn't scale to full distribution.

Answer 3

B — Vector Clocks ⚠️ honest but incomplete

Vector clocks track causality. Each write carries a version vector — one counter per replica. When two writes have no causal relationship (neither "happened before" the other), you know you have a conflict. Both versions surface to the application layer, which resolves them.

Amazon's internal Dynamo used this. It's correct and transparent. The problem: conflict resolution is now your job. You write merge logic for every data type. For shopping carts it's manageable (union the items). For arbitrary documents? It breaks down fast.

Vector clocks tell you *where* conflicts are. They don't tell you *how* to fix them.


Answer 4

A — Last-Write-Wins (LWW) ❌ dangerous

Simple. Brutally simple. Highest timestamp wins. The loser's change vanishes — no error, no notification, just gone.

Cassandra defaults to this. For append-only workloads (metrics, event logs, telemetry) it's totally fine. For anything a human edited? It's silent data loss with a timestamp attached. The treacherous part: users never get an error. Their work just stops being there.

Fitting a Large Document into an LLM



Your LLM has 128K tokens.

Your document has 150K words.

Something has to give. What do you do?

- A** Chunk the document into fixed-size pieces and embed each one — retrieve the top-k at query time.
- B** Use a sliding window — process the document in overlapping chunks, stitch the outputs together.
- C** Summarize each section progressively — feed the running summary forward as context.
- D** Truncate to the most recent tokens and hope the answer is near the end.


Answer 1

✓ CORRECT ANSWER

WHY A WINS: RAG WITH CHUNKED EMBEDDINGS

System Design Labs

✓ Split the document into overlapping chunks, embed each one, store in a vector DB, retrieve top-k at query time.



Why A wins (RAG with chunked embeddings):

You split the document into overlapping chunks (~500 tokens, ~10–20% overlap), embed each one, store in a vector DB (Pinecone, pgvector, Qdrant), and retrieve top-k at query time.

This scales to any document size. Your LLM context at inference only sees the retrieved chunks — staying well inside the context window. Latency stays low because retrieval is a fast ANN lookup, not a sequential scan.

The catch: **chunking boundaries destroy semantic coherence**. If a clause starts on page 12 and resolves on page 13, a hard chunk boundary between them means neither chunk retrieves well for that question. This is why overlap and chunk size tuning matter so much. Most teams underestimate this.

Production fix: chunk at sentence or paragraph boundaries, not character counts. Hybrid retrieval (keyword + semantic) handles the edge cases.

Answer 2

Why B is the trap answer (sliding window):

Sliding window looks elegant — process the document in overlapping passes, accumulate the answer. But it doesn't actually fix the context window problem. You still have to process every chunk sequentially, and the answer to a question spanning chunk 1 and chunk 47 gets lost. You're doing $O(n)$ LLM calls per query. At 150K words and a 3s/call LLM, that's minutes per query. Nobody ships this in production for QA on large documents.

It works fine for **summarization** tasks where the output from chunk N feeds into chunk N+1. For retrieval? Wrong tool.

Answer 3

Why C sometimes works (progressive summarization):

For summarizing a 200-page doc? Solid. You map-reduce the document — summarize sections, then summarize the summaries. GPT-4o does this well.

For answering a specific question about clause 38(b) on page 97? It doesn't. The intermediate summary will have compressed away the specific detail you needed. Lossy compression applied to information retrieval is a category error.

Use progressive summarization when you want a high-level understanding. Don't use it when specific facts need to survive to the end of the pipeline.

Answer 4


Why D is obviously wrong (and also surprisingly common):

Truncating to the last N tokens is the default behavior of most LLM API wrappers when you exceed the context limit and don't handle it explicitly. Many teams ship this by accident.


The problem beyond the obvious data loss: LLMs have a "**lost in the middle**" **failure mode** (Liu et al., 2023). Performance on retrieval tasks is highest when the answer is at the start or end of the context window — and degrades significantly when it's buried in the middle. So even if you fit everything in context, position matters.

Truncation compounds this — you're not just losing data, you're losing data unpredictably.

Preventing Concurrent Balance Overspend



Two users read \$200 from the same wallet at the same time. Both spend \$150. Wallet ends up at -\$100. How do you stop this?



A

Pessimistic locking — `SELECT FOR UPDATE` on the wallet row. One transaction blocks until the other commits.

B

Optimistic locking — read the row with a version number, only write if version hasn't changed since you read it. Retry on conflict.

C

MVCC — let both reads see a consistent snapshot, rely on the database to detect write conflicts at commit time.

D

Serializable isolation — set the transaction isolation level to `SERIALIZABLE` and let the database handle it.

You have a payment system. Two users try to spend from the same wallet balance at the same time.

Both read \$200. Both want to spend \$150. Both see enough balance. Both write the deduction.

The wallet is now at -\$100. How do you stop this?

A

Pessimistic locking — `SELECT FOR UPDATE` on the wallet row. One transaction blocks until the other commits.

B

Optimistic locking — read the row with a version number, only write if version hasn't changed since you read it. Retry on conflict.

C

MVCC — let both reads see a consistent snapshot, rely on the database to detect write conflicts at commit time.

D

Serializable isolation — set the transaction isolation level to `SERIALIZABLE` and let the database handle it.


Answer 1

✓ CORRECT ANSWER

B Why B wins (optimistic locking):



You read the wallet row, grab the version number. When you write, you include a **WHERE version = :read_version** condition. If another transaction already updated the row, your **WHERE** matches **zero rows** — conflict detected, retry.



Update Query

```

UPDATE wallets
SET balance = balance - 150, version = version + 1
WHERE id = :wallet_id AND version = :read_version;
  
```



If rows_affected = 0 → conflict → retry.



No locks held during the read. At **10K TPS** with **low conflict rates**, this is significantly faster than pessimistic locking — you're only paying for retry cost on actual conflicts, not lock acquisition on every read.



The catch:

under **HIGH contention** (same hot wallet hit repeatedly), **retry storms** eat you alive.



Redis INCRBY with atomic decrement

OR



Queue in front of the wallet update

Why B wins (optimistic locking):

You read the wallet row, grab the version number. When you write, you include a **WHERE version = :read_version** condition. If another transaction already updated the row, your **WHERE** matches zero rows — conflict detected, retry.

UPDATE wallets

SET balance = balance - 150, version = version + 1

WHERE id = :wallet_id AND version = :read_version;

If rows_affected = 0 → conflict → retry. No locks held during the read. At 10K TPS with low conflict rates, this is significantly faster than pessimistic locking — you're only paying for retry cost on actual conflicts, not lock acquisition on every read.

The catch: under HIGH contention (same hot wallet hit repeatedly), retry storms eat you alive. Hot wallets need Redis INCRBY with atomic decrement, or a queue in front of the wallet update.

Answer 2

Why A is the trap answer (pessimistic locking):

SELECT FOR UPDATE works. It's correct. But it holds a row-level lock for the entire transaction — every concurrent write queues up.

At 10K TPS on a popular wallet, you've serialized all writes to a single queue. Throughput collapses. Lock wait timeouts cascade — failed transactions retry, adding more contention. You've built a self-reinforcing bottleneck.

Use pessimistic locking for low-concurrency paths (admin ops, batch jobs) where conflict probability is near 100%. Not for high-throughput payment writes.

Answer 3

Why C is the silent failure (MVCC + default isolation):

This is the dangerous one. PostgreSQL defaults to READ COMMITTED. MVCC gives each transaction a consistent snapshot — but at READ COMMITTED, that snapshot refreshes per statement, not per transaction.

Two concurrent transactions both read \$200, both pass the "sufficient funds" check, both commit. No conflict detected. **Double spend in production.**

MVCC prevents dirty reads. It does NOT prevent lost updates at READ COMMITTED. Most engineers assume MVCC = safe from concurrent writes. It doesn't.

Answer 4


Why D is almost always the wrong default (SERIALIZABLE):

Theoretically correct — the database guarantees transactions behave as if they ran serially. No anomalies.

The problem: PostgreSQL's SSI tracks read-write dependencies across transactions. Under high concurrency, serialization failures spike — transactions abort and retry even when there was no actual conflict. Real throughput penalty at scale.

Right for complex financial workflows where correctness > throughput (end-of-day reconciliation). Wrong default for a high-throughput per-transaction API.

Offloading Heavy Work Off the Main Thread



Your React dashboard freezes for 4 seconds every time a user uploads a CSV.

Not a network freeze. A *browser* freeze. Tab unresponsive. Animations stuck. The "Page Unresponsive" dialog popping up on Chrome.

Here's what's happening in prod:

- User drops a 80MB CSV (200k rows, 40 columns) into the upload zone
- Your code: `Papa.parse(file, { complete: ... })` on the main thread
- Parsing + validation + schema mapping = 3.8s of synchronous JS
- During those 3.8s: scroll dies, the loading spinner stops spinning, input lag goes infinite
- Your INP score on the perf dashboard is now red. Lighthouse is screaming.

Product wants this fixed before the enterprise demo on Friday. Four options on the table:

- A** `requestIdleCallback` + time-slicing — chunk the parse into 5ms slices, yield between rows, let the browser breathe.

B Web Worker — move the entire parse + validation pipeline into `worker.js`, communicate via `postMessage`, main thread stays free.

C `SharedArrayBuffer` + `Atomics` — share the file buffer between main thread and 4 workers, parallelize the parse across cores, zero-copy.

D Compile the parser to WebAssembly — drop the JS parser, run a Rust CSV parser through WASM for near-native speed.

Answer 1


✓ CORRECT ANSWER

Why B Wins: Web Worker



Move the heavy CSV parse + validation off the main thread.

Keep the UI smooth. Always.



Why B wins:

The actual problem isn't *speed*. It's *where the work runs*. The main thread is the same thread that handles scroll, input, layout, paint, and animation. Anything synchronous on it = jank. Period.

A Web Worker is a separate OS-level thread with its own event loop. You ship the file (transferable via `postMessage` with a `Transferable` so it's zero-copy), parse it over there, send back the parsed rows. Main thread stays at 60fps the entire time. INP drops from 3800ms to ~80ms. Done.

This is what Figma does for its rendering pipeline, what Google Sheets does for formula recalc, what Excalidraw does for its scene processing, what VS Code Web does for syntax highlighting. Every serious browser-based product runs heavy compute in a worker. It is the boring, correct, ten-year-old answer.

Answer 2

Why A is wrong (the "clever" trap):

requestIdleCallback + time-slicing keeps the UI *technically* responsive but stretches the parse from 3.8s to ~6s of wall-clock time, because you're constantly yielding. Worse: idle callbacks are deprioritized — under load, they don't fire at all. React 18's concurrent rendering uses time-slicing for *rendering*, not for blocking IO-shaped work. Wrong tool.

Answer 3

Why C is wrong (the SENIOR ENGINEER TRAP):

SharedArrayBuffer + Atomics is the answer that sounds smart in a design review. "We'll parallelize across cores! Zero-copy! Lock-free!" Then you discover:

- It requires COOP + COEP headers site-wide, which break Stripe Elements, YouTube embeds, Intercom, and half your analytics scripts.
- A single 80MB CSV parse is *embarrassingly serializable* — splitting rows across 4 cores saves you ~2s while costing you weeks of cross-origin isolation hell.
- You now own a custom thread pool in JavaScript. In 2026.

SAB is the right answer for *exactly two* problems: real-time audio/video processing (where you need lock-free ring buffers between threads at 48kHz), and shared simulation state (think multiplayer game engines in the browser). A one-shot CSV upload is neither.

Answer 4

Why D is wrong:

WASM makes the parser *faster*. It doesn't move it *off the main thread*. A WASM parser running synchronously on the main thread still freezes the UI — just for 1.2s instead of 3.8s. You'd still need to put it inside a worker. So D isn't an answer; it's a possible optimization *inside* B. And honestly, Papa Parse in a worker is fast enough that you don't need it.

Moving from Batch to Real-Time Streaming



Your data team just got a new SLA: surface fraud signals within 500ms of a transaction.

Right now you're running nightly Spark batch jobs. The business wants "real-time." Your team knows Spark. Someone already opened a PR adding Spark Structured Streaming.

The transaction volume: 8,000 events/sec peak. You're on AWS. The fraud model runs in Python. The output feeds a DynamoDB table the API reads from.

You need to redesign the pipeline. What do you pick?

- A) Kafka Streams — event-by-event processing, stateful operators, sub-10ms latency. Lives inside your app JVM.
- B) Apache Flink — true streaming engine, exactly-once semantics, built for high-throughput stateful processing.
- C) Spark Structured Streaming — micro-batch under the hood, 100ms–5s windows, same API your team already knows.
- D) Keep the batch job, drop the window to 1 minute — "near real-time" at zero migration cost.

Answer 1

✓ CORRECT ANSWER


System Design Labs

WHY B WINS: APACHE FLINK

 **SUB-50ms**
end-to-end
latency

 **8K**
events/sec
handled easily

 **EXACTLY-ONCE**
no double-counting
on failures

 **WATERMARKS**
handle out-of-order
events cleanly


TRUE STREAMING, NOT BATCH

Flink is a true streaming engine — event-by-event. Every transaction triggers processing the moment it arrives.



LOW LATENCY

Latency: sub-50ms end-to-end with proper tuning. Meets the 500ms SLA with headroom.



HIGH THROUGHPUT

At 8K events/sec, Flink doesn't blink. Built for high-throughput stateful workloads.



EXACTLY-ONCE SEMANTICS

State checkpoints + two-phase commits ensure your fraud signals aren't double-counted when a node fails.



HANDLES OUT-OF-ORDER EVENTS

Flink's watermark model handles out-of-order events cleanly — crucial with clock skew from distributed payment processors.



PYTHON FRAUD MODEL INTEGRATION

Integrate via PyFlink (embedded Python operators) or a sidecar Python microservice over gRPC.



OPERATIONAL REALITY

Operational overhead is real. More components. More knobs. Requires skilled operators. But the SLA demands it.



BUILT FOR THIS PROBLEM

You're not reaching for Flink prematurely. The business asked for 500ms and gave you 8K events/sec. That's exactly the problem Flink was built for.

Why B wins (Apache Flink):

Flink is a true streaming engine — event-by-event, not batch. Every transaction triggers processing the moment it arrives. Latency: sub-50ms end-to-end with proper tuning. At 8K events/sec, Flink doesn't blink. Exactly-once semantics mean your fraud signals aren't double-counted when a node fails. The Python fraud model integrates via PyFlink or a sidecar microservice. Flink's watermark model handles out-of-order events cleanly — which matters when transactions arrive from distributed payment processors with clock skew.

Operational overhead is real. But the SLA demands it. You're not reaching for Flink prematurely here — the business asked for 500ms and gave you 8K events/sec. That's exactly the problem Flink was built for.

Answer 2

Why A is close but not the cleanest fit (Kafka Streams):

Kafka Streams is event-by-event too — sub-10ms achievable. The constraint: it runs inside your application JVM. Your fraud model is Python. That means Kafka Streams handles the stream topology, but the model scoring lives in a separate Python service, adding a network hop. Flink handles that boundary more cleanly with PyFlink. Kafka Streams is the right call if the fraud logic is stateful JVM-native aggregation. Wrong default when the model is a Python service.

Answer 3

Why C is the senior engineer trap (Spark Structured Streaming):


The PR is already open. The API looks familiar. But Structured Streaming is micro-batch — it collects events for a trigger interval (practically 100ms–5s under production load), then processes them as a mini-batch. You'll tune it, it'll pass staging, and it'll breach SLA the first time peak traffic hits. Mistaking "Streaming" in the name for actual event-by-event processing is the trap.

Answer 4

Why D is eliminated immediately (batch, 1-minute windows):

Minimum latency = the window size. At 60 seconds you're 120x over SLA before the first line of fraud logic runs. "Near real-time" is not a SLA.

Designing Memory for an AI Agent



Your AI agent remembered the user's name.

Then it forgot what it was doing.

Here's the setup:

User asks the agent: book the cheapest flight to NYC, search hotels under \$150/night, then compare total trip cost.

By step 3, the agent calls the LLM with 8,000 tokens of raw conversation history — and still answers as if it's turn 1.

You need a memory architecture before this ships. Which one do you pick?

A) In-context window only — full conversation stays in the system prompt. Simple. Breaks at ~15 turns or 8K tokens, whichever comes first.

B) Vector memory store — embed past turns, retrieve the top-k by semantic similarity at query time. Works great until "NYC flight" pulls a memory about a past NYC trip instead of the current task.

C) Episodic memory with summarization — compress old turns into structured event summaries, inject the relevant ones per request. More complex to build. Much harder to confuse.

D) Redis session state — structured key-value store, explicit agent reads/writes. Deterministic. Requires the agent to know what to store and when.

Answer 1

✓ CORRECT ANSWER



C — Episodic memory with summarization

System Design Labs

Right answer for task-oriented agents.




→ Keep the last N turns in context (short-term)



→ Compress older turns into structured event summaries (episodic)



→ Inject only relevant summaries per new request



Why this works at scale

- Stable context window
- Harder to confuse
- Degrades gracefully
- Used by Anthropic, LangChain Memory v2, and OpenAI Assistants API



A In-context window only — full conversation stays in the system prompt. Simple. Breaks at ~15 turns or 8K tokens, whichever comes first.



B Vector memory store — embed past turns, retrieve the top-k by semantic similarity at query time. Works great until "NYC flight" pulls a memory about a past NYC trip instead of the current task.



C Episodic memory with summarization — compress old turns into structured event summaries, inject the relevant ones per request. More complex to build. Much harder to confuse.



D Redis session state — structured key-value store, explicit agent reads/writes. Deterministic. Requires the agent to know what to store and when.

C — Episodic memory with summarization

Right answer for task-oriented agents. The pattern:

- Keep the last N turns in context (short-term)
- Compress older turns into structured event summaries (episodic)
- Inject only relevant summaries per new request

It's how production agents at Anthropic, LangChain Memory v2, and OpenAI Assistants API all converge. More engineering upfront — but the only pattern that degrades gracefully at scale.

Answer 2

A — In-context window only

Collapses fast. At ~15 turns or 8K tokens you're either truncating early context or hitting the model's limit. Works for demos. Fails in production agents with multi-step tasks. The "just increase context length" argument ignores cost and the lost-in-the-middle problem — LLMs reliably ignore content buried in the middle of long prompts.

Answer 3

B — Vector memory store


The retrieval problem is nastier than it looks. Semantic similarity isn't the same as task relevance. In a multi-step booking task, embedding "find hotels under \$150" might retrieve a memory from a different user session — or a past conversation about a different city entirely. Great for knowledge retrieval. Unreliable for procedural task state.

Answer 4

D — Redis session state

Solid and deterministic. The catch: the agent has to explicitly decide what to write to state and when to read it. That's a non-trivial design problem — you're building a working memory protocol on top of your tool calls. Most teams underestimate the schema design this requires.

Reliable CDN Cache Invalidation on Deploy



Your CDN saved you from explaining a 400ms load time to your VP.

Then your deploy shipped. And nobody got the fix.

Here's what happened:

Your edge nodes were still serving the old JS bundle — TTL hadn't expired. You triggered a cache purge manually. But you hit the wrong environment. Prod kept serving the stale version for 6 more hours.

You need a CDN configuration that handles invalidation reliably under deploy pressure. Which one do you pick?


- A) Anycast routing + stale-while-revalidate + event-driven purge via API on deploy
- B) Unicast routing + aggressive TTLs + manual cache purge via dashboard
- C) Anycast routing + long TTLs + tag-based invalidation via API
- D) Anycast routing + short TTLs + no invalidation strategy (let caches expire naturally)

Answer 1

✓ CORRECT ANSWER

A Anycast + stale-while-revalidate + event-driven purge via API on deploy

System Design Labs



A — Anycast + stale-while-revalidate + event-driven purge — the right answer

Anycast advertises the same IP from multiple PoPs. BGP routes each user to the closest edge automatically. Tokyo user hits Tokyo — not Virginia.

Stale-while-revalidate: serve the cached version immediately, fetch a fresh copy in the background. Users never wait on a cache miss to resolve.

Event-driven purge: your CI/CD pipeline fires a Cloudflare or CloudFront API call on every deploy, targeting exactly the changed assets. No manual steps. No wrong-env clicks. No 2 AM mistakes.

Answer 2

C — Anycast + long TTLs + tag-based invalidation — solid, but has a trap

Tag-based invalidation (Cloudflare Cache-Tag, Fastly surrogate keys) lets you group objects and purge entire logical groups in one API call. New product images? Invalidate tag:product-images. Clean and surgical.

The trap: long TTLs work perfectly when your deploy pipeline is robust. But if your invalidation API call fails silently during an incident — when you're moving the fastest — you're serving the wrong version for hours with no TTL fallback to save you.

Option A's stale-while-revalidate degrades more gracefully when your pipeline gets messy.

Answer 3

B — Unicast + aggressive TTLs + manual purge

Unicast means one edge region. Your Tokyo users are routing to Virginia. You've deployed a CDN and you're routing past it.

Manual dashboard purges during a deploy: someone will forget, click the wrong environment, or skip it entirely at 2 AM. Aggressive TTLs shrink the damage window — they don't make invalidation reliable.

Answer 4


D — Short TTLs, no invalidation

"Short TTLs = always fresh" is the most expensive misconception in CDN config.

Every TTL expiry is a potential cache stampede. 200k RPS at peak, TTL expires on a hot endpoint, every PoP fires a simultaneous request to your origin. You go from 0 to 200k RPS in 2 seconds.

And you have zero control for hotfixes. You just wait for the TTLs to drain.

Syncing Offline Edits Without Data Loss



Your team just shipped "offline mode" for your field-service app.

Technicians work in basements, tunnels, plant floors — connectivity is unreliable.

The demo looked great. [Navigator.onLine](#) said false, the app kept working, the sync button pulsed green. You shipped to 400 users.

Then the incident reports started.

Technician in Munich finishes a repair job offline, syncs when she gets signal. Her updates are gone — overwritten by a colleague who edited the same record online 12 minutes earlier. Last-write-wins. Her write lost.

Technician in São Paulo opens the app, goes offline, edits three assets. Comes back online. App throws an unhandled promise rejection and crashes. IndexedDB schema was on version 2. The update shipped version 3. The migration never ran because he'd never opened the app while online.

You're now asked to actually fix offline-first — not demo it.

Here's the setup:

- 400 field technicians, avg offline window of 40 minutes

- Write conflicts happen ~3x per day, always on the same 8 "hot" assets
- Sync currently runs on reconnect via a single bulk POST
- IndexedDB is being used but schema migrations are undocumented
- You need conflict resolution, migration safety, and sync reliability

What's your architecture?

A) Move everything to localStorage + a manual JSON diff on sync. Simpler API, deterministic schema versioning — no IndexedDB migration headaches.

B) Keep IndexedDB, add a vector-clock field to every record. On sync, compare clocks — if diverged, surface a merge UI. Let the technician decide. Never throw away a write.


C) Wrap IndexedDB in a versioned schema migration layer (like Dexie.js). Add a per-record updated_at + device_id composite key. Last-write-wins, but last-write is now deterministic and auditable.

D) Replace client-side storage with a CRDT-based sync engine (Automerge or Yjs). Operations are commutative and associative — merge is always valid, conflicts are structurally impossible.

Answer 1

✓ CORRECT ANSWER

Why D wins: CRDT Engine — Automerge / Yjs



Bottom line: If you need offline + multi-writer + zero data loss, CRDTs are the only approach with a mathematical guarantee.

Why D wins (CRDT engine — Automerge / Yjs):

CRDTs (Conflict-free Replicated Data Types) are the only approach where merge is mathematically guaranteed to produce the same result regardless of the order operations arrive. They're commutative, associative, and idempotent.

Two technicians can edit the same asset record offline for 6 hours, sync in any order — and the final state is always deterministic. No merge UI. No "who wins" logic. No data loss.

Linear uses Automerge. Figma uses a custom CRDT for multiplayer. Notion runs a similar model for their block store. This isn't overengineering — it's what teams reach after trying everything else.

Tradeoff: CRDT payloads are larger (they carry operation history). Yjs is lighter than Automerge for most cases. Server-side merge logic needs to understand CRDT ops, not just overwrite rows.

Answer 2

Why B is the correct second choice (Vector Clocks + Merge UI):

Vector clocks detect *exactly* when a conflict occurred — which device wrote what and in what order. When diverged, you surface a merge UI and let the technician decide.

Right answer when humans must stay in the loop. A technician marking an asset "inspected" vs another marking it "needs repair" — that's not a conflict an algorithm should auto-resolve.

Tradeoff: every conflict requires user action. At 3/day it's manageable. At 300/day it's a support nightmare.

Answer 3

Why C is good but not enough (Versioned IndexedDB + LWW):

Dexie.js makes IndexedDB migrations genuinely manageable. LWW with a deterministic composite key (updated_at + device_id) is production-grade.

But LWW throws away real writes. The Munich technician *always loses* to anyone who touched the same record online in the interim. In a field-service app where "inspected" vs "not inspected" has regulatory consequences — that's data loss you can't accept.

LWW works when the cost of losing a write is low (a UI preference). It breaks when every write has operational meaning.

Answer 4


Why A fails (localStorage + JSON diff):

localStorage is hard-capped at ~5MB. At 200 records with attachment metadata, you'll hit it — silently, with no useful error in older Chrome. No indexing. No range queries. In private/incognito mode, it's scoped per-tab and wiped on close.

JSON diff breaks on nested arrays, partial field updates, deletions. You'd end up building a mini-CRDT by hand without the guarantees.

Works in a weekend prototype. Corrupts silently in production.

Scaling Full-Text Search



Your search box returns results in 80ms.

Then you add 50M more documents. Now it's 4 seconds.

Your DBA says "add an index." Your search engineer says "that's not how this works."

They're both right — about different things.

Here's the setup:

Platform: SaaS product, 100M documents

Search fields: title, description, tags, partial phrases

Current stack: PostgreSQL full-text search

Current perf: p95 = 4.2s

Target: sub-200ms at p99

Product is escalating. What do you do?

- A** Migrate to Elasticsearch — inverted index, purpose-built for full-text, sub-100ms at scale.

B Add GIN indexes on your PostgreSQL tsvector columns — no new infra, meaningful performance gain.

C Cache the top 1,000 most common queries in Redis — free win for the majority of traffic.

D Move to Typesense or Meilisearch — simpler ops than Elasticsearch, built for search.

Answer 1

✓ CORRECT ANSWER


**ANSWER: A – Elasticsearch with inverted indexes**

System Design Labs

WHY A WINS

Elasticsearch uses inverted indexes. Queries do fast posting-list intersections instead of table scans. Scales to 100M+ docs with horizontal sharding and parallel query fan-out for sub-100ms at p99.

Target:
sub-200ms at p99

**A**

Migrate to Elasticsearch — inverted index, purpose-built for full-text, sub-100ms at scale.

B

Add GIN indexes on your PostgreSQL tsvector columns — no new infra, meaningful performance gain.

C

Cache the top 1,000 most common queries in Redis — free win for the majority of traffic.

D

Move to Typesense or Meilisearch — simpler ops than Elasticsearch, built for search.

Why A wins (Elasticsearch):

Elasticsearch is built on an inverted index. When you index a document, Elasticsearch tokenizes the text and writes posting lists: "database" → [doc_1, doc_5, doc_23...], "index" → [doc_2, doc_5, doc_18...]. A query for "database index" does a fast posting-list intersection — not a table scan.

At 100M documents, that intersection still runs in milliseconds because posting lists are stored in delta-compressed sorted form, Lucene segments are immutable (no row locking, no MVCC overhead), horizontal sharding is native (queries fan out in parallel), and BM25 scoring runs during retrieval — not after it.

Result: sub-100ms at p99, even at 100M+ docs, with proper shard sizing. Netflix, Uber, and GitHub all run some form of this at their search layer.

The tradeoffs you accept: new infrastructure to operate (index lifecycle, shard rebalancing), write amplification from segment merging, and eventual consistency — by default, newly indexed docs aren't searchable for ~1 second.

Answer 2

Why B is the trap answer (GIN indexes on PostgreSQL):

GIN indexes on tsvector are genuinely good. For 1M–5M documents, this is the right call — no new infra, no ops burden, real improvement.

At 100M documents, you hit three walls: index size balloons (GIN is large — easily 40–80GB), PostgreSQL full-text has no native horizontal read scaling, and BM25 relevance scoring is bolted on, not native. You get from 4.2s down to maybe 500ms. You need 200ms. Close, but the wrong tool at this scale.

GIN is what you reach for before you need a dedicated search engine, not instead of one.

Answer 3

Why C misses the point (Redis cache):

Caching top queries is a real optimization — but it's not a search architecture. Your top 1,000 queries might cover 40% of traffic. The other 60% hits the slow path. Worse: every document update invalidates your cache logic. You've added complexity without touching the tail latency product is actually complaining about.

C is a band-aid. The p99 problem is still there the moment any user types something that isn't in the cache.


Answer 4

Why D is tempting but limited (Typesense / Meilisearch):

Typesense and Meilisearch are excellent — simpler ops than Elasticsearch, great developer experience, built-in typo tolerance. For 10M–30M documents, D is a legitimate answer that wins on simplicity.

At 100M documents with strict p99 requirements and complex relevance tuning, you start hitting their limits: distributed indexing is less mature, fine-grained scoring customization is constrained, and operational tooling at scale is thinner than Elasticsearch's ecosystem. For a smaller product, D is the right call. At 100M with a p99 SLA, Elasticsearch is where you end up.

Enforcing Structured Output from an LLM



Your production LLM agent just returned this JSON to your order processing service:

```
{
  "action": "refund",
  "amount": "fifty dollars",
  "order_id": null,
  "confidence": "pretty high"
}
```

Your downstream service crashes. The retry hits the same model. Same broken output. The refund never fires — but the user got a confirmation email.

You need your agent to return valid, typed, structured output — every time. What do you do?


- A)** Prompt-engineer harder — add "Always return valid JSON with these exact fields" to your system prompt and document the schema inline.
- B)** Use structured outputs / function calling (OpenAI, Bedrock tool use, Gemini response schema) — constrain the model at the API level to return a typed schema.
- C)** Post-process with a validation layer — parse the output, run JSON Schema or Pydantic validation, retry with corrective context if it fails (max 2 retries).
- D)** Add a second LLM as a judge — pass the first model's output to a smaller, faster model that scores and flags invalid responses before they reach your service.

Answer 1 ✓ CORRECT ANSWER



B WINS: Structured Outputs / API-Level Schema Enforcement

The only option that eliminates the problem at the source.



Why B wins (Structured outputs / API-level schema enforcement):

This is the only option that eliminates the problem at the source. When you use OpenAI's `response_format: { type: "json_schema" }`, Bedrock's tool use / converse API, or Gemini's `response_schema`, the model's token sampling is *constrained* — it cannot produce output that violates the schema. Not "it's less likely." Cannot.

Under the hood, these APIs use constrained decoding: the model's output probability distribution is masked so invalid tokens are zeroed out at each step. The result is guaranteed schema conformance on every call. No retries. No parsing. No "it worked 98% of the time."

In production: use Pydantic (Python) or Zod (TypeScript) to define your schema, pass it to the API, deserialize directly into your typed model. Your downstream service never sees a string where it expects an int.

Answer 2

Why A is wishful thinking (Prompt engineering):

"Always return valid JSON" is not a contract. It's a suggestion.

LLMs are probabilistic text predictors. Under load, edge-case prompts, or slightly different phrasing, they drift. You'll get valid JSON 97% of the time and a markdown code block the other 3%. At 10k requests/day that's 300 failures. At 100k, it's 3,000.

Prompt engineering is a first step when you're prototyping. It is not an output validation strategy for production. Teams that rely on this eventually write 50-line system prompts trying to enumerate every failure mode. The model still escapes.

Answer 3

Why C is correct but incomplete (Post-process + retry loop):

Validation layer + retry with corrective context is a legitimate, widely-used pattern — especially when you can't use structured outputs (e.g., open-source models like Llama, Mistral, or your own fine-tune via Ollama/vLLM).

The implementation: parse → validate with Pydantic/JSON Schema → on failure, re-inject the original prompt + the bad output + a correction instruction ("Your previous output had these errors: [list]. Try again.") → retry, max 2 attempts → fail hard and surface to human review.

But: this is your fallback, not your primary. Every retry is latency + cost. Under high load, retry storms can spike your API spend 3–4x. Corrective context also doesn't help when the model fundamentally can't produce what you're asking — you'll just retry into the same failure.

Use C when B isn't available. Use B when it is....

Answer 4

Why D is a trap (LLM-as-judge for validation):


LLM-as-judge is a real eval pattern — it's used to score *quality*, tone, safety, and correctness in offline evaluation pipelines and red-teaming.

For *structural* output validation in a hot path? It's expensive, slow (~200–500ms extra per request), and unreliable for schema enforcement. A judge model can tell you "this response seems wrong" — it cannot guarantee a valid typed object. You've added latency and cost without actually solving the problem.

LLM-as-judge belongs in your eval harness (nightly runs, CI) — not in your request/response loop for data integrity.

The worse failure mode: teams add a judge *in addition to* bad structured prompting, feel safer, and ship it. Both models can hallucinate in sync. The judge agrees the broken output looks fine. Silent corruption ships to prod.

Incrementally Strangling a Monolith



Your monolith is 6 years old.

500k lines of code. One deploy every 3 weeks. One bad migration takes down the whole thing.

Leadership says: "rewrite it." You've seen what happens when teams do that. 18 months. Budget overruns. Half the features never make it back. The new system launches and nobody trusts it.

There's a better way. You don't rewrite — you strangle.

Here's your system:

- OrderService (monolith) handles: order creation, fulfillment, invoicing, returns
- 40 engineers. 3 teams. All deploying to the same codebase.
- Returns processing is a bottleneck. The business wants it extracted first.
- You can't freeze feature work during migration.

You need to extract Returns into a standalone service without a big-bang rewrite.

What's your migration strategy?

- A) Strangler Fig** — route /returns traffic to a new service via a proxy/facade, keep the monolith intact until the new service is proven. Deprecate monolith code after.
- B) Branch by Abstraction** — introduce an interface inside the monolith, build the new implementation behind it, flip the switch when ready, then move it out.
- C) Big Bang Rewrite** — freeze the monolith, rebuild everything in microservices in parallel, cut over when done.
- D) Database-First Migration** — extract the returns database tables into a separate schema first, build the service around it, then reroute traffic.

Answer 1


✓ CORRECT ANSWER

A**Strangler Fig – the right way to migrate**

Route `/returns` to a new service via a proxy, keep the monolith intact, and deprecate old code only after the new service proves itself.

**Why A wins**

- ✓ Migrate live, no freeze
- ✓ Rollback at every step
- ✓ Proven at scale (Netflix, Amazon)

**Why A wins:**

You introduce a routing proxy in front of the monolith. `/returns` traffic routes to the new service. The monolith handles everything else. You run both in parallel — no freeze, no big-bang. Migrate one endpoint at a time, each step independently rollback-able. Only deprecate monolith code after the new service has proven itself under real load. Netflix and Amazon have run versions of this at scale.

Answer 2**Why B is close but different (Branch by Abstraction):**

Real pattern — but it lives *inside* the codebase. You wrap the logic behind an interface, build the new implementation behind a flag, then flip it. It's a refactoring tool, not a deployment strategy. Often used *inside* a Strangler Fig migration to prepare internal structure. But on its own it doesn't get you to an independently deployed service.

Answer 3

Why C is wrong (Big Bang Rewrite):


The monolith doesn't freeze while you rebuild. Business keeps shipping features. Your parallel rewrite immediately diverges from reality. "Done" keeps moving. The cutover is a single high-stakes event with no incremental validation. Joel Spolsky called this "the single worst strategic mistake any software company can make" — because every hack in the old codebase was a real bug fix you've now forgotten.

Answer 4

Why D is more dangerous than it looks (Database-First):

Extracting the DB tables first sounds logical but breaks the monolith's ORM immediately. You've created a distributed data problem *before* building the distributed system to handle it. Data migration is the *last* step — extract service logic first, keep data co-located temporarily, then migrate with dual-write. Doing it in reverse order compounds the risk at every step.

How an Agent Chooses Which Tools to Call



Your AI agent just got a user message: "Book me a flight to Dubai next Friday."

The LLM has access to 12 tools: search_flights, get_user_preferences, check_calendar, book_flight, send_confirmation, get_weather...

How does the agent decide which tools to call, in what order, and when to stop?


- A) ReAct loop** — model reasons step-by-step, emits a "thought" then picks one tool at a time, observes output, repeats until it self-decides it's done
- B) Parallel tool calling** — model emits ALL required tool calls in a single response, executes them concurrently, feeds all results back in one context update
- C) Forced function schema** — you lock the model into a strict JSON schema per turn; it can't produce free text, only structured tool calls you defined
- D) Planner-executor split** — a lightweight planner LLM creates a tool call DAG upfront, a separate executor runs the graph, results flow back to planner only at checkpoints

Answer 1

✓ CORRECT ANSWER

A**ReAct loop** ✓ Correct for most cases

ReAct (Reason + Act) is the dominant production pattern for a reason.




The model emits a "Thought:" explaining what it's doing, then an "Action:" with a single tool call, then observes the result before deciding what's next.

Why it works:

- Each step is observable and debuggable
- The model can bail out, retry, or change direction after each tool result
- Works even with tools that have side effects — you can gate dangerous calls
- Supported natively by OpenAI, Anthropic, LangChain, LlamaIndex out of the box

The trap:

ReAct is sequential. If you need `get_user_preferences` AND `check_calendar` AND `search_flights` all before `book_flight`, you're making 3 round trips when you could make 1.

**A — ReAct loop** Correct for most cases

ReAct (Reason + Act) is the dominant production pattern for a reason.

The model emits a "Thought:" explaining what it's doing, then an "Action:" with a single tool call, then observes the result before deciding what's next.

Why it works:

- Each step is observable and debuggable
- The model can bail out, retry, or change direction after each tool result
- Works even with tools that have side effects — you can gate dangerous calls
- Supported natively by OpenAI, Anthropic, LangChain, LlamaIndex out of the box

The trap: ReAct is sequential. If you need `get_user_preferences` AND `check_calendar` AND `search_flights` all before `book_flight`, you're making 3 round trips when you could make 1.

Answer 2

B — Parallel tool calling Right when dependencies allow it

Modern LLMs (GPT-4o, Claude 3.5+) can emit multiple tool calls in a single response.

If the tools are independent — run them in parallel. Massive latency win.

Example: `get_user_preferences` + `check_calendar` have no dependency on each other. Call both at once. Then feed both results back before calling `search_flights`.

The trap: not all tools are independent. If Tool B needs the output of Tool A, parallel calling breaks your context. The model doesn't "know" B failed because A hadn't run yet — it just gets a bad result and may hallucinate a fix.

Production pattern: use parallel calling for the fan-out phase, ReAct for the sequential decision phase.

Answer 3

C — Forced function schema The senior engineer trap

Looks like control. Actually limits you.

When you force strict JSON schema mode, you lose the model's ability to reason before calling.

This breaks down fast when:

- The user's intent is ambiguous (should the agent call `search_flights` or `clarify_intent`?)
- A tool fails and the model needs to decide whether to retry or give up
- You need conditional logic between steps

Use schema enforcement for output validation, not tool dispatch. Validate what comes back from the model — don't handcuff the model going in....

Answer 4

D — Planner-executor split 🏗️ Right at scale, overkill for most

This is the architecture behind production agentic pipelines like Devin, AutoGPT-style systems, and enterprise agent frameworks.

The planner builds a DAG of tool calls upfront. The executor runs it. The planner only re-engages at checkpoints or on failure.

Benefits:


- Parallelism by default (the DAG has clear dependency edges)
- Cost-efficient (the powerful planner model runs less)
- Auditable — you can inspect the plan before running it

The trap: the planner has to get the full plan right before execution starts. If the world changes mid-run (API returns unexpected data, user updates their request), you need replanning. That's a whole second problem.

Most teams reach for this too early. Start with ReAct. Add parallel tool calling when latency hurts.

Only build planner-executor when you have complex multi-step workflows with 10+ tools and real parallelism requirements....

Controlling Data-Warehouse Query Costs



Your data team just opened a \$4,200 BigQuery bill.

For a single month. One analyst. 12 queries.

The queries weren't wrong. They weren't inefficient SQL. They were reasonable analytics queries — "give me last 30 days of events for customer X." The problem was that every single one scanned the full 3.2 TB table. No partition pruning. No cost control. Just full scans, every time.

This is the most expensive silent bug in data engineering. You write a query. It looks fast. It returns results. And every run quietly eats through terabytes you're paying per-byte to scan.

The fix is partition strategy — but picking the wrong one doesn't just fail to help, it actively makes things worse.

Here's the setup:

You're running a 3.2 TB events table on BigQuery. 18 months of data. Ingested daily. Analyst queries almost always filter on two things: a date range ("last 30 days") and a customer_id ("for customer X").

Which partition + clustering strategy do you pick?

A) Partition by ingestion date — date range queries only scan the relevant day-partitions.

- B)** Partition by `customer_id` ranges — customer queries only scan the relevant ID bucket.
- C)** No partitioning — use Redshift `DISTKEY` on `customer_id` + `SORTKEY` on `event_time`.
- D)** Partition by ingestion date + cluster on `customer_id` — date pruning cuts the partition, clustering prunes within it.

Answer 1

✓ CORRECT ANSWER

D Why D wins: Partition by ingestion date + cluster on customer_id



BigQuery
Partition + Cluster



1. Partition by date

BigQuery skips every partition outside the date range before scanning a single row.

For a "last 30 days" query on 18 months of data, you immediately eliminate ~94% of the table.



2. Cluster by customer_id

Within each scanned partition, BigQuery sorts the rows by customer_id into storage blocks.

A filter on customer_id now skips 90–95% of the remaining rows.




Combined impact

A "last 30 days for customer X" query goes from scanning 3.2 TB → ~12 GB.

Bill drops from \$4,200 → ~\$70/month.

WITHOUT PARTITIONING + CLUSTERING

18 MONTHS OF DATA
3.2 TB



Query: last 30 days
for customer X

Scans 3.2 TB
\$4,200/month

WITH PARTITION BY DATE + CLUSTER BY CUSTOMER_ID

18 MONTHS OF DATA
3.2 TB



Partition pruning (date)



Scans ~12 GB
\$70/month

✓ Same queries.
Same data.
Right strategy.



This pattern works across every major warehouse:



BigQuery
Partition + Cluster



Snowflake
Micro-partitioning +
Cluster Keys



Redshift
Partition Key +
SORTKEY Compound

Why D wins:

Your analyst queries always filter on two columns: a date range and a customer_id. A single-key strategy protects one and ignores the other.

Here's what D does in BigQuery:

- **Partition by date** → BigQuery skips every partition outside the date range before scanning a single row. For a "last 30 days" query on 18 months of data, you immediately eliminate ~94% of the table.
- **Cluster by customer_id** → Within each scanned partition, BigQuery sorts the rows by customer_id into storage blocks. A filter on customer_id now skips 90–95% of the remaining rows.

Combined: a "last 30 days for customer X" query goes from scanning 3.2 TB → ~12 GB. Bill drops from \$4,200 → ~\$70/month. Same queries. Same data. Right strategy.

This pattern works across every major warehouse: BigQuery (partition + cluster), Snowflake (micro-partitioning + cluster keys), Redshift (partition key + SORTKEY compound).

Answer 2

Why A is incomplete (date partition only):

Date partition alone is a real improvement — date range filters prune the irrelevant partitions immediately. Your bill probably drops 60–70%.

But your analyst queries also filter by `customer_id`. Inside each date partition, all customer rows are interleaved across storage blocks. BigQuery has to read the entire partition to find customer X's rows. You've cut the scan from 3.2 TB → ~500 GB (the last 30 days), but you're still scanning every customer inside that window.

A is the obvious first step. D is A taken all the way.

Answer 3

Why B is the senior engineer trap (`customer_id` range partitioning):

This one gets smart people. Your analysts always filter by `customer_id` — why not partition on that?

Because real `customer_id` distributions are never uniform. Your enterprise clients (IDs 1–10,000) generate 80% of all events. Every analytics query against active customers hits the same partition. No pruning benefit — the entire interesting data is in one bucket.

Worse: your queries also filter by date. BigQuery now has to scan the full customer partition to find rows from the last 30 days. You've traded a full table scan for a full-partition scan on your hottest partition. Same cost, different framing.

Hot partitions on high-traffic ID ranges is the most common expensive mistake in data warehouse design. Looks right in theory. Never works in production.

Answer 4

Why C fails at this scale (DISTKEY + SORTKEY, no partitioning):

DISTKEY distributes rows across Redshift nodes by `customer_id` hash — optimizes joins. SORTKEY on `event_time` sorts rows on disk within each slice — helps date range scans read contiguous blocks.


Both are real optimizations. Neither replaces partition pruning at 3.2 TB.

Without partitioning, every query triggers a full table scan across all nodes. If your DISTKEY distribution is skewed — which it will be when enterprise customers dominate — one Redshift node holds 60–70% of your data. Every expensive query hammers that node.

You've turned a cost problem into a latency + cost problem.

DISTKEY + SORTKEY is solid for join optimization and smaller tables. It's not a substitute for partition pruning at terabyte scale.

Serving ML Models at High Throughput



Your ML team trained a model that hits 94% accuracy in the notebook.

Then you deploy it.

Peak traffic: 3,000 inference requests/second.

P99 latency shoots to 4.2 seconds.

GPU utilization: 23%.

The model works. The serving layer is the bottleneck.

Here's the setup:

- PyTorch model, ~6B params
- Single A100 GPU, 80GB VRAM
- FastAPI wrapper calling `model.predict()` one request at a time
- No batching, FP32 weights, no quantization

Users are hitting timeouts. GPUs are sitting mostly idle. Your infra bill is climbing.

What do you fix first?


- A)** Enable dynamic batching — buffer incoming requests, group into batches, process in one forward pass.
- B)** Quantize the model to INT8 — reduce weight precision from 32-bit to 8-bit, shrink memory footprint and speed up inference.
- C)** Switch to tensor parallelism — split the model across multiple GPUs, distribute the compute.
- D)** Add a request queue + async workers — decouple HTTP receiving from model inference, process jobs in background.

Answer 1 ✓ CORRECT ANSWER


**Answer: A — Dynamic Batching**

System Design Labs

Batch requests for a few milliseconds, run one forward pass, and keep the GPU fed.

**BEFORE:** No Batching (One Request at a Time)

Symptoms: 3K RPS, P99 = 4.2s, GPU Util = 23%
GPU is starving. High overhead, low throughput.

**AFTER:** Dynamic Batching (Many Requests per Forward Pass)

Result: Higher throughput, lower P99 latency, happier users, and dramatically better GPU utilization.

**Why A Wins**

- ✓ The GPU is underutilized (23%) → it's not the compute, it's the request pattern.
- ✓ Dynamic batching turns many small, expensive calls into fewer, efficient batch calls.
- ✓ Massively improves throughput and latency without changing the model.
- ✓ This is the first lever to pull for low GPU util + high RPS.

Why Others Don't Fix This First

- B** Quantization helps throughput, but doesn't fix low GPU utilization.
- C** Tensor parallelism helps if one GPU is maxed out. It's not.
- D** Queues/async help resilience, not GPU utilization or latency.

Answer: A — Dynamic Batching ✓

Here's the breakdown:

Why A wins (Dynamic Batching):

The smoking gun is the combination: 3K RPS, P99 = 4.2s, GPU utilization = 23%.

GPU utilization at 23% tells you the GPU is starving. It's not overloaded — it's bored. Every request is arriving one at a time, paying the full GPU kernel launch overhead, then idling while the next request trickles in. GPUs are built for parallelism. They want 256 inputs to process simultaneously, not 1.

Dynamic batching buffers requests for a few milliseconds (say, 50ms), collects 32–128 of them, and fires a single forward pass. That same GPU that was at 23% utilization can jump to 85–90%. Latency improves because batch throughput dramatically outpaces sequential per-request calls, even with the small buffer wait.

This is exactly how NVIDIA Triton Inference Server, TorchServe, and vLLM (for LLMs) work by default. Batching is the first thing you reach for when GPU util is low and RPS is high.

Answer 2

Why B is the trap answer (INT8 Quantization):

Quantization is legitimate and valuable — but it solves a *different* problem. It reduces model size (~4x smaller) and speeds up individual forward passes. If you were memory-constrained or couldn't fit the model at all, quantization is critical. But this model fits comfortably in 80GB VRAM and the GPU isn't saturated.

INT8 quantization also introduces precision loss you need to validate carefully, and the speedup on a modern A100 for a 6B param model in isolation is modest — maybe 1.5–2x. That doesn't fix 3K RPS with sequential inference. You'll still have sequential execution, just slightly faster per call.

The right order: batch first, quantize second (for memory or cost, not latency emergencies).

Answer 3

Why C is wrong here (Tensor Parallelism):

Tensor parallelism splits model layers across multiple GPUs using NVLink or interconnects. It reduces per-request latency for *very large* models that don't fit on one device, or when you need sub-100ms latency on a single huge model.

But this is a 6B param model on an 80GB A100 — it fits. Tensor parallelism adds coordination overhead and complexity, and it doesn't help when the bottleneck is sequential request processing. You'd also need 2+ GPUs, which inflates cost before you've even tried the cheap fix.


Answer 4

Why D is wrong (Async Queue + Workers):

An async queue decouples HTTP from inference and prevents timeout failures — that's a reliability win. But it doesn't improve GPU utilization or actual throughput. You're still processing one request at a time, just asynchronously. Users now get job IDs instead of timeouts, but the wall-clock wait is the same or worse.

Queuing without batching is adding a lobby to a slow restaurant. It hides the congestion — it doesn't fix it.

Isolating Noisy Tenants in Multi-Tenant SaaS



You're building a B2B SaaS product. 50 enterprise customers. Each one wants their data isolated. Some are on free plans. A few are paying \$50k/year and demanding SLA guarantees.

Your current setup:

- One database. One schema. A `tenant_id` column on every table.
- One app server handling all traffic.
- A free-tier customer running a badly-written bulk export just hammered your DB for 40 seconds. A paying enterprise customer's checkout flow timed out.

Your investors are not happy. Neither is that enterprise customer.

The engineering question: how do you isolate tenants without rebuilding the whole product?

- A) Keep one shared DB — add row-level security + query budgets per tenant to enforce limits.
- B) Schema-per-tenant — every customer gets their own schema in the same Postgres instance, migrations run per-schema.
- C) Database-per-tenant (silo model) — each enterprise customer gets a dedicated DB. Free tier stays pooled.

D) Middleware bridge — route requests to tenant-specific DB clusters based on a tenant registry, free tier stays on shared pool.

Answer 1

✓ CORRECT ANSWER

Why D wins: Middleware Bridge — Hybrid Model

Route each tenant to the right data home. Isolate where it matters. Scale where it's cheap.

REQUEST FLOW



- Free / Small (Shared)
- Mid-tier (Schema-per-tenant)
- Enterprise (DB-per-tenant)

THE PAYOFF

- Noisy neighbor is contained.**
A runaway free-tier export can only impact other free-tier tenants.
- Enterprise customers can't touch each other. Period.**
Dedicated DBs provide strong isolation and SLA guarantees.
- Grow tenants up the isolation ladder without code changes.**
Move a tenant between tiers by updating the registry only.
- GDPR deletion is a drop, not a cascade.**
One DB drop for enterprise tenants. Clean, fast, verifiable.



How it works

At request time, the app server asks the Tenant Registry where a tenant lives. It gets back a connection config (host, port, db, schema) and connects. The application code stays the same for every tenant.



Why D wins (Middleware bridge — hybrid model):

This is the pattern Salesforce, HubSpot, and every mature B2B SaaS converges on eventually.

The architecture:

- Free/small tenants → shared pool DB (cost-efficient, acceptable risk)
- Mid-tier tenants → schema-per-tenant in a shared Postgres cluster
- Enterprise tenants → dedicated DB cluster (full isolation, SLA-able)

A tenant registry (Redis or a lightweight service) maps `tenant_id` → connection config at request time. Your app server doesn't care which model a tenant uses — it asks the registry and gets a connection string back.

The payoff:

- Noisy neighbor is contained. A runaway free-tier export can only hurt other free-tier tenants.
- Enterprise customers can't touch each other. Period.

- You can migrate a tenant up the isolation ladder as they grow — without touching app code.
- GDPR deletion is one DB drop for enterprise tenants, not a cascade delete across shared tables.

Answer 2

Why A fails (Row-level security + query budgets):

RLS works. Query budgets can work. But they share the same DB process. A tenant that opens 500 connections, triggers a table lock, or fills the WAL — that bleeds into every other tenant on the instance. You can throttle CPU, but you can't throttle I/O blast radius once it's in the buffer pool.

This is the band-aid answer. Fine at 5 tenants. Falls apart at 50.

Answer 3

Why B is dangerous (Schema-per-tenant):

Schema-per-tenant sounds clean. It is, until:

→ 300 tenants need a migration. That's 300 sequential DDL operations on one Postgres instance. Migration windows become a multi-hour risk event.

→ At 500 tenants you're burning connection slots faster than you can PgBounce your way out.

→ One long-running query in schema_tenant_247 still holds locks on shared system tables. Noisy neighbor hasn't gone away — you just can't see it as clearly.

Schema-per-tenant is a local maximum. You hit Postgres scaling walls before you hit the point where it pays off.

Answer 4

Why C is too expensive too early (Full silo):

The silo model is the correct isolation model for enterprise customers. But applying it to *all* tenants — including a free tier with 1,000 accounts — means 1,000 Postgres instances. That's \$40k/month in RDS costs before you've made a dollar.

Full silo for everyone is the right endgame for the top 5% of your customer base. Applying it everywhere is overengineering.

Versioning an API Without Breaking Clients



Your API just shipped a breaking change.

/users now returns fullName instead of first_name + last_name. 3 mobile clients broke. 1 partner integration went down. Your on-call is not happy.

You had a versioning strategy. It just wasn't the right one.

There are 4 ways to version an API. Here's what actually happens when you pick each one in production:

A — URL path versioning (/v1/users, /v2/users)

Simple. Explicit. Every request makes the version visible in logs and caches. But now you're maintaining 2 full route trees. A bugfix in the business logic layer has to be patched in both. Teams quietly let v1 rot.

B — Header versioning (API-Version: 2)

Clean URLs. Version negotiation in the transport layer, not the path. Harder to test in a browser, invisible in logs unless you instrument for it, and clients forget to send the header — defaulting to whatever your server decides "latest" means.

C — Query param versioning (/users?version=2)

Fast to implement. Zero client SDK changes. Cache-unfriendly — every CDN layer treats ?version=1 and ?version=2 as separate cache keys. Works until you have 40 endpoints and version drift becomes untrackable.

D — Content negotiation (Accept: application/vnd.api.v2+json)

The REST-purist approach. Semantically correct — you're asking for a *representation*, not a route. Almost nobody implements it right. Client library support is inconsistent. One wrong Accept header and you get a 406.

Answer 1

✓ CORRECT ANSWER

System Design Labs

D Content negotiation: correct spec, wrong reality


`Accept: application/vnd.api.v2+json`

is what the HTTP spec intended.


You're asking for a specific representation of a resource, not a different route.

Semantically, it's the most correct.

HOW CONTENT NEGOTIATION SHOULD WORK



THE REALITY IN PRODUCTION



Bottom line:

The most correct on paper, but the least reliable in practice. Great in theory. Painful in production.

`Accept: application/vnd.api.v2+json`

→ 406 Not Acceptable

D — Content negotiation: correct spec, wrong reality

`Accept: application/vnd.api.v2+json` is what the HTTP spec intended. You're asking for a specific *representation* of a resource, not a different route. Semantically, it's the most correct.

In practice: Accept header parsing is inconsistent across client libraries. Middleware strips headers silently. A misconfigured client gets a 406 Not Acceptable with no obvious error message pointing to the version field. GitHub tried content negotiation early in their API and moved to URL versioning. That tells you something.

Answer 2

A — URL path versioning: the pragmatic default (not wrong, but has a cost)

This is what Stripe, Twilio, and GitHub use. Not because it's theoretically pure — it isn't — but because it's operationally obvious.

Every proxy, log aggregator, and CDN understands it without configuration. `/v1/` vs `/v2/` is visible everywhere without instrumentation.

The cost is real: you're running parallel route trees, and v1 maintenance debt compounds fast. A bugfix in shared business logic has to be patched in both versions. Stripe's answer was to pin clients to a version at the API key level and deprecate slowly over years. Most teams don't have that discipline — so v1 quietly rots while v2 gets all the attention.

Answer 3

B — Header versioning: good theory, rough edges

Clean architecture. The URL represents a resource, the header represents the negotiation. This is how internal microservices often version — when you control both client and server, it works well.

The problem: invisible in browser tabs, impossible to share as a curl command without extra flags, and "default to latest" on a missing header is a footgun. In year two, a legacy client quietly drops the header. Your server defaults it to v2. The client breaks. Nobody knows why — it's not in the URL, it's not in the logs unless you specifically instrumented for it.


Answer 4

C — Query params: the quick fix that stays forever

Fast to implement. Zero client SDK changes. Works immediately.

Cache invalidation gets messy at scale — every CDN layer treats `?version=1` and `?version=2` as separate cache keys, so your cache hit rate tanks. The bigger problem is version drift: with no forcing function, clients scatter across v1/v2/v3 with no deprecation pressure. Good for internal tools where you control all the clients. Gets expensive for public APIs.

Zero-Downtime Schema Migrations



Your migration ran fine in staging.

Then you ran it in production.

The app went down.

Not because the SQL was wrong. Because you ran it on a live table with 40 million rows while 8 services were actively writing to it.

Your setup:

→ PostgreSQL. users table. 40M rows. Active writes from 8 services.

→ Product request: split full_name into first_name + last_name.

→ You have a 2-hour maintenance window tonight.

The engineering question: how do you ship this without downtime?

A) Run ALTER TABLE to drop full_name and add first_name + last_name in a single migration during the maintenance window.

B) Add first_name + last_name as nullable columns first → backfill → update all services to write to both → drop full_name only after everything is migrated.

- C)** Create a new users_v2 table with the target schema → dual-write to both tables → flip the read pointer → drain the old table.
- D)** Add a DB view that aliases full_name as first_name || ' ' || last_name → let each service migrate off it at its own pace.

Answer 1 ✓ CORRECT ANSWER


Why **B** is correct (expand/contract)

System Design Labs

✓ Two phases. Expand safely. Then contract when everything is ready. Every step is reversible.

PHASE 1
EXPAND

- 1 Add first_name and last_name as nullable columns.
- 2 Backfill in small batches (rate-limited, no row lock escalation).
- 3 Update services to write to both columns.




✓ WHY THIS IS SAFE

- Adding a nullable column with no default in PostgreSQL is a metadata-only operation — near-instant, no lock.
- Backfill in small batches keeps the system responsive.
- Old column remains source of truth until you're ready.
- Every step is independently reversible.

PHASE 2
CONTRACT

- 1 Once all services write to the new columns and reads are fully migrated...
- 2 Drop full_name.
- 3 You control the timing.



✓ RESULT

- full_name is dropped only when it's safe.
- You control the exact timing.
- Minimal risk. Maximum control.

Slower than A? Yes.
3am pages? Zero.

Why **B** is correct

Two phases.

Phase 1 — Expand: add first_name and last_name as nullable columns. On PostgreSQL, adding a nullable column with no default is a metadata-only operation — near-instant, no lock. Backfill in small batches (rate-limited, no row lock escalation). Update services to write to both columns simultaneously.

Phase 2 — Contract: once every service is writing to the new columns and reads are fully migrated, drop full_name. You control the timing.

The key property: every step is independently reversible. If something breaks after Phase 1, you stop — the old column still exists, old services still work.

Slower than A? Yes. 3am pages? Zero.

Answer 2

Why A will get you paged

ALTER TABLE on a live table acquires an ACCESS EXCLUSIVE lock in PostgreSQL. It blocks every read and write for the duration. On a 40M row table with active traffic, that lock holds for minutes — not seconds. Your 8 services time out. The queue backs up. Alerts fire.

CONCURRENTLY helps for index creation. It doesn't help for column drops or renames. The maintenance window doesn't save you — the lock is the problem, not the timing.

Answer 3

Why C is overkill

Shadow table + dual-write works. But it requires a dual-write layer in every service, an atomic read-pointer flip, and a drain/cleanup phase for the old table.

Worth it for genuine structural rewrites — new primary key, re-partitioning, changing storage engines. For a column split on a single table? That's 3 weeks of engineering for a 3-day problem.


Answer 4

Why D defers the problem

A view aliasing full_name buys time. Useful as a transition layer inside Option B. But PostgreSQL write-through views only work on simple single-table views with no joins or aggregates. You still need to backfill the real columns eventually.

It defers the migration. It doesn't replace it.

Handling Embedding Drift in RAG



You built a RAG pipeline. Works great in dev.

6 months later, your users complain: "The search results are garbage."

You haven't changed a line of code.

Here's what happened:

Your product evolved. New features, new docs, new support tickets. The data drifted — but your embedding index didn't.

Now you're serving a 400GB FAISS index that was last rebuilt in January. Your chunks are stale. Your nearest-neighbor results point to deprecated docs. Your LLM is confidently hallucinating from outdated context.

You need to fix this. 4 engineers each propose a solution:

A) Scheduled full rebuild

Every Sunday, re-embed the entire corpus from scratch. Replace the index atomically. Slow (4h+ at scale), expensive, but always fresh.

B) Incremental upserts + soft delete

On every document change, re-embed only the affected chunks. Mark deleted chunks as tombstoned. Keep a version field on each vector. Index size grows over time; compact quarterly.

C) Embedding version registry + hot swap

Track which embedding model version produced each vector. When the model drifts (fine-tuned or upgraded), invalidate the mismatched vectors and rebuild only those. Two indexes run in parallel during migration. Route traffic by model version.

D) Approximate staleness detection

Run a nightly job that samples 1% of your corpus, re-embeds it, and measures cosine distance against the stored vector. If drift exceeds a threshold, trigger a full rebuild. Otherwise, skip it. Cheap monitoring, reactive rebuilds.

Real constraint: your corpus is 50M chunks. Full rebuild = 4 hours + ~\$800 in embedding API cost. You deploy model updates every 6 weeks.

Answer 1

✓ CORRECT ANSWER




C is the right answer — Embedding Version Registry + Hot Swap

This is the only solution that handles the real problem: your embedding model changed.



C tracks which model produced which vector and routes queries to the correct index version during migration. Two indexes in parallel during transition = zero downtime, zero mixed-space comparisons.



PRODUCTION WIN



Gradual Migration
Migrate 10% of traffic to the new index, validate retrieval quality, then increase confidently.



Zero Downtime
Two indexes run in parallel. No big bang. No 4-hour Sunday rebuild blocking your team.



Safe Validation
Compare results, monitor quality, roll back instantly if needed.



Correctness
No mixed embedding spaces. Cosine similarity math stays valid.

WHY A, B, AND D FAIL



Scheduled Full Rebuild
Still creates a mixed index during migration if the model changed. Big bang, expensive, risky.



Incremental Upserts + Soft Delete
Index ends up with vectors from multiple models. Cosine similarity breaks across embedding spaces.



Approximate Staleness Detection
Detects drift reactively, but rebuild is still a big bang and causes downtime.



The Real Problem
Different embedding models produce different vector spaces. You can't compare them.

C is the right answer — Embedding Version Registry + Hot Swap

This is the only solution that handles the *real* problem: your embedding model changed. When you fine-tune or upgrade your model, Option A and B still serve a mixed index — some vectors from model v1, some from v2. Your cosine similarity math breaks because you're comparing vectors from different embedding spaces. They're not comparable. C is the only option that tracks *which model produced which vector* and routes queries to the correct index version during migration. Two indexes in parallel during transition = zero downtime, zero mixed-space comparisons.

The production win: you can migrate 10% of traffic to the new index, validate retrieval quality, then cut over. No big bang. No 4-hour Sunday rebuild blocking your team.

Answer 2

B is the senior engineer trap

Incremental upserts sound smart — you're only re-embedding what changed. But here's the silent killer: model drift. You update your embedding model every 6 weeks. Now your index has vectors from 4 different model versions. Chunk A was embedded with text-embedding-ada-002. Chunk B with your fine-tuned v2. Nearest neighbor search across a mixed-space index returns garbage — not because the data is stale, but because the *geometry* is inconsistent. B handles data drift. It doesn't handle model drift. At 6-week update cycles, model drift is your actual problem.

Answer 3

A is brute-force ops debt


\$800 every Sunday whether you need it or not. At 6-week model updates, most rebuilds are wasted spend on data that hasn't meaningfully changed. It works. It's just expensive and blunt. Fine for a startup at 1M chunks. Painful at 50M.

Answer 4

D is monitoring, not a fix

Staleness detection is genuinely useful — but as a signal, not a strategy. The 1% sample tells you *that* you have drift, not *which* vectors are affected or *why*. You still need a rebuild strategy. D is a trigger, not an architecture. Combine it with C for a complete system.

Parallel Agents with Shared State



You built an agent. It works.

Now you need 5 of them running in parallel, sharing state, and handing off work to each other.

Your pipeline breaks on the first real workload.

Here's the setup:

You're building a research agent system. A user asks a complex question. You need to:

- Fan out to 3 specialized sub-agents simultaneously
- One agent might spawn 2 more based on what it finds
- They all write back to shared context
- A final agent synthesizes everything

Classic multi-agent orchestration. You have 4 options for how agents coordinate.

A) Centralized Orchestrator — one controller agent dispatches tasks, collects results, manages shared state. Agents are dumb workers.

B) Decentralized Peer Handoff — each agent decides who gets the task next. No central controller. Agents communicate directly.

C) Shared Message Queue + Blackboard — all agents read/write to a shared blackboard. Coordination happens through state, not calls.


D) Hierarchical Nesting — orchestrator spawns sub-orchestrators. Each sub-tree is self-coordinating. Recursive decomposition of the problem.

Answer 1

✓ CORRECT ANSWER

**CENTRALIZED ORCHESTRATOR**

One controller agent owns the plan, dispatches tasks, collects results, manages state.

**WHY A WINS****SINGLE SOURCE OF TRUTH**

The orchestrator owns the plan and state. Everyone else executes.

**EASY TO DEBUG & OBSERVE**

Failures are centralized. Logs, traces, and metrics are complete.

**RESILIENT BY DESIGN**

Sub-agents can fail or restart without corrupting global state.

**DETERMINISTIC & REPLAYABLE**

You can replay an orchestrator run from its state log end-to-end.

**PRODUCTION PROVEN**

LangGraph, CrewAI, and production systems default to this model.

**KEY TAKEAWAY**

Centralize control, distribute execution. The orchestrator plans, dispatches, collects, and synthesizes. Sub-agents stay simple, stateless, and reliable.

- ✓ Clear ownership
- ✓ Better reliability
- ✓ Easier debugging
- ✓ Scales with confidence

Answer: A — Centralized Orchestrator

Here's the full breakdown — and where each pattern actually belongs:

Why A wins (Centralized Orchestrator):

One agent owns the plan. It dispatches, collects, retries, and synthesizes. Every sub-agent is stateless and dumb on purpose — they receive a task, return a result, done.

This isn't a limitation. It's a feature.

When something fails, you know exactly where to look: the orchestrator. It holds the retry logic, the timeout budgets, the fallback paths. Sub-agents can die and restart without corrupting global state because the orchestrator is the only one maintaining it.

LangGraph, CrewAI, and every production-grade agentic framework defaults to this model for a reason. It's debuggable, observable, and deterministic. You can replay an orchestrator run from its state log. Try doing that with peer-to-peer handoffs.

Answer 2

Why B is the trap (Decentralized Peer Handoff):

This sounds elegant in theory. Agents self-organize, no bottleneck at the top.

The problem: who resolves conflicts? If Agent A hands off to Agent B, and Agent B decides the task actually belongs to Agent C, which also hands it back to Agent A — you have a cycle with no circuit breaker.

Peer handoff works in closed, well-defined pipelines (like a Git commit hook chain). It breaks the moment you have conditional routing, error recovery, or parallel fan-out.

Answer 3

Why C is partial credit (Shared Blackboard):

A shared blackboard is real and useful — but as a *data layer*, not a coordination mechanism.

Agents writing to and reading from shared state works for passing context. What it can't do cleanly is orchestrate *who runs next, when, and with what inputs*. You end up building an implicit orchestrator via polling — same thing, no visibility into it.

Use a blackboard for state sharing. Use an orchestrator for task routing. They're complementary, not alternatives.


Answer 4

Why D fits a specific problem (Hierarchical Nesting):

Hierarchical orchestration is the right model when the problem itself is hierarchically decomposable — like a pipeline where "summarize AI papers" fans out to 10 sub-tasks, each of which fans out further.

The catch: recursion adds latency and debugging complexity at every layer. If your problem doesn't naturally decompose into clean sub-trees, you're adding indirection for no gain. Default to flat centralized orchestration first. Reach for hierarchical only when the problem demands it.

Tracking Long-Running Async Jobs



Your background job ran for 4 minutes and nobody knows if it finished.

That's not a job queue problem. That's a missing design problem.

Long-running jobs break every assumption you built for synchronous APIs. Your load balancer times out after 30s. Your mobile client doesn't know whether to retry. Your retry logic re-runs a job that already half-completed.

Here's the real scenario:

You're processing a video upload. The job takes 2–8 minutes. Millions of users.

What do you expose to the client?

- A) Polling endpoint — client hits `/jobs/:id/status` every 5s until done
- B) Webhook — job fires a POST to client's callback URL on completion
- C) SSE / WebSocket — server pushes progress updates in real time
- D) Synchronous wait — keep the HTTP connection open until the job finishes

One scales to millions without coupling your infrastructure to client uptime.

The others have hard production failure modes most teams don't discover until 3 AM.

The deeper problem isn't transport — it's these 4 things nobody gets right the first time:

→ **Idempotency.** Every job must be safe to re-run. If your retry logic can double-charge, double-send, or double-process — you don't have retries, you have bugs waiting.

→ **Progress granularity.** "0% → 100%" is useless for a 6-minute job. You need intermediate states: queued, processing, transcoding, uploading, complete. Clients need something to show users.

→ **Timeout vs failure.** A job that stops responding isn't the same as a job that failed. Dead workers, OOM kills, spot instance evictions — your queue needs a heartbeat or a visibility timeout, not just a try/catch.


→ **Deduplication.** The client will retry. Your queue will redeliver. You need a dedup key scoped to the original request — not the job run.

Answer 1

✓ CORRECT ANSWER

A) POLLING — CORRECT FOR MOST CASES

System Design Labs

**Why Polling Wins**

Client-controlled, stateless, scales independently. The server doesn't care if the client disconnects, retries, or crashes — the job runs and the status endpoint just answers queries.

5s polling intervals on a job that takes 2–8 minutes is trivially cheap.

LinkedIn, YouTube, and S3 multipart uploads all use polling for async job status.

THE KEY RULES**Stable & Idempotent Job IDs**

Make job IDs stable and idempotent so retries return the same job.

**TTL on Job Records**

Set a TTL on job records so you don't accumulate state forever.

**Exponential Backoff**

Use exponential backoff instead of fixed intervals to reduce load.

**Stateless Status Endpoint**

The status endpoint should be fast, cacheable, and read-only.

A) Polling — Correct for most cases

Client-controlled, stateless, scales independently. The server doesn't care if the client disconnects, retries, or crashes — the job runs and the status endpoint just answers queries. 5s polling intervals on a job that takes 2–8 minutes is trivially cheap. The key rules: make job IDs stable and idempotent, set a TTL on job records so you don't accumulate state forever, and use exponential backoff not fixed intervals. LinkedIn, YouTube, and S3 multipart uploads all use polling for async job status.

Answer 2**B) Webhook — Right idea, wrong default**

Webhooks are great for server-to-server flows where the receiver has a stable HTTPS endpoint. They fall apart for mobile clients (no public URL), in environments with NAT/firewall, and when the receiver is down at delivery time. You'd need a retry queue, delivery guarantees, and signature verification just to make it reliable. Webhooks work well as a supplementary delivery mechanism for platform integrations — not as the primary client notification path for end users.

Answer 3

C) SSE / WebSocket — Expensive for this use case


Real-time push is fantastic for chat, live dashboards, and collaborative editing. For a job that takes 2–8 minutes and is triggered once? You're holding an open connection, burning a file descriptor, and adding connection-management complexity for maybe 3 meaningful state changes. SSE makes sense when you genuinely need sub-second updates or continuous streaming. For async job progress, polling is cheaper and simpler.

Answer 4

D) Synchronous wait — Production antipattern

Keeping the HTTP connection open for 2–8 minutes kills your load balancer (most timeout at 30–60s), ties up a server thread or process, and gives the client no recovery path if the connection drops mid-job. This is how systems get "stuck" — jobs complete but the client never hears about it because the connection dropped at minute 3. Never block on long-running work. Always return immediately with a job ID.

Consistency for Payment-Confirmation Reads



Your write traffic just hit 50k writes/sec. Reads are 10x that.

You split your Postgres into one primary + two replicas. Reads go to replicas. Now you're shipping a payment confirmation flow — and a replica returned stale data 200ms after a write. The user saw "payment pending" when it already succeeded.

Here's the setup:

- Primary handles all writes
- Two read replicas, async replication
- Payment confirmation reads from a replica
- Replication lag: 80–300ms under load

You need reads-after-write consistency on payment flows without tanking write throughput. What do you do?

- A** Switch to synchronous replication — primary waits for at least one replica to confirm before ack'ing the write.

B Use semi-sync replication — primary waits for exactly one replica to write to its relay log, but not flush to disk.

C Route payment confirmation reads back to the primary, keep everything else on replicas.

D Add a short sleep (500ms) after every payment write before reading — wait out the lag.

Answer 1


✓ CORRECT ANSWER

Why C wins (primary reads for critical paths):

System Design Labs



Payment confirmations are a tiny slice of your read traffic. You're not routing ALL reads to primary — just the reads where stale data causes **visible user harm**.



This pattern is called **read-your-writes consistency**. Stripe, Shopify, every serious fintech does this. You add a flag or middleware: "if this is a consistency-sensitive flow, read from primary." Everything else keeps hitting replicas.



Minimal impact

Payment confirmations are maybe 0.1% of your reads. Primary read traffic increases slightly.



Surgical precision

Only critical flows read from primary. Everything else stays on replicas.



No new failure modes

No topology change. No new components. No added complexity.



Maximum consistency where it matters

Users never see "payment pending" after a successful payment again.



Bottom line: Read from primary for consistency-sensitive flows. Fast, safe, and battle-tested.

Why C wins (primary reads for critical paths):

Payment confirmations are a tiny slice of your read traffic. You're not routing ALL reads to primary — just the reads where stale data causes visible user harm.

This pattern is called **read-your-writes consistency**. Stripe, Shopify, every serious fintech does this. You add a flag or middleware: "if this is a consistency-sensitive flow, read from primary." Everything else keeps hitting replicas.

Primary read traffic increases slightly — but a "just confirmed payment" flow is maybe 0.1% of your reads. Negligible. No topology change, no new failure modes, surgical precision.

Answer 2

Why A is tempting but the wrong trade (synchronous replication):

Postgres `synchronous_commit = on` — primary waits for a replica to confirm before acking the write. Zero lag. Reads always consistent.

The cost: every write now takes `primary_write_time` + network roundtrip to replica. 1–5ms normally. Under load with a flaky replica? Your write latency is hostage to your slowest replica.

If that sync replica goes down → you either halt writes or fall back to async and break your guarantee. That choice gets made at 3am under pressure.

Sync replication is a serious architectural commitment for the whole DB. Routing payment reads to primary solves the same problem, surgically.

Answer 3

Why B is the sneaky wrong answer (semi-sync replication):

Semi-sync (`rpl_semi_sync` in MySQL) — primary waits for one replica to write to its relay log. Sounds like a middle ground. It's not.

The replica acknowledged receipt of the bytes — not that the data is applied and readable. The relay log is the *incoming queue*, not the applied state. You can still read stale data immediately after a write.

Semi-sync protects you from **data loss on primary failure**. It does NOT protect you from replication lag on reads. Different problem entirely.

Senior engineers conflate "replica has the data" with "replica has applied the data." They're not the same thing.

Answer 4

Why D will destroy you at scale (sleep-based lag compensation):


Sounds pragmatic. Works for about 2 weeks.

Replication lag isn't a fixed constant. It's 50ms now, 300ms during a write spike, 2s during a vacuum. Hardcode a sleep and you've built a system that's either too slow (`sleep > lag`) or still broken (`lag > sleep`).

Every payment confirmation path is now artificially 500ms slower. And when the sleep isn't enough during a spike — you're back to the original bug, with worse latency.

Sleep-based consistency is always a smell.

Observability for LLM Agent Pipelines



It's 2am. Your LLM agent is on fire and every log says HTTP 200.

You shipped an AI-powered support agent 6 weeks ago. It calls tools, chains 4-5 prompts per request, hits vector search, calls back into your API. Users are complaining. Answers are wrong. Token bills are 3x what they were last week.

You open your logs. Everything is green. You open your traces. You see one span called POST /chat that took 14 seconds. That's it. That's your entire visibility.

Here's the setup:

NestJS API → OpenAI (GPT-4o + GPT-4o-mini) + Anthropic (Claude Sonnet fallback)

~8,000 agent runs/day, each chains 3-7 LLM calls + 2-4 tool calls

Token spend jumped from \$180/day to \$540/day this week, no idea which flow caused it

Eval suite passes 94% in CI, prod users are hitting hallucinations you can't reproduce

On-call engineer literally cannot answer "why is this run slow / expensive / wrong?"

You have one sprint to fix the visibility problem before the next incident. What do you instrument first?

- A)** Structured trace IDs through every LLM + tool call, span-level token accounting, latency per hop — LangSmith, Helicone, or Arize as the backbone.
- B)** LLM-as-judge eval pipeline running on a sample of prod traffic — assertion-based checks, ragas metrics, catch quality regressions before users do.
- C)** Cost dashboards per model/agent/user with hard budget alerts — stop the token bleed first, everything else can wait.
- D)** Custom middleware that logs every prompt in/out, model, latency, tokens to your own store — own your data, no vendor lock-in, query it however you want.

Answer 1

✓ CORRECT ANSWER



Why A wins:

You can't fix what you can't see.

System Design Labs



Slow?

Which of the 7 LLM calls was slow.



Expensive?

Which model on which step burned the tokens.



Wrong?

Which retrieval returned garbage that poisoned the next prompt.

→ **None of those questions have answers without a trace.** LLM traces aren't HTTP traces — you need span-level attributes: model name, prompt tokens, completion tokens, cost, tool name, retrieval doc IDs, and the actual prompt/response payload. That's what LangSmith, Helicone, Arize, and Langfuse give you out of the box.

Before: One 14s span. Zero answers.

POST /chat

200

14.02s

14.02s



You have no idea what happened inside that 14 seconds.

After: The 14-second span opens into 7 sub-spans.

POST /chat

200

14.02s

Agent Run

14.02s



1. Classify Intent (GPT-4o-mini)

1.21s | 320 in | 88 out | \$0.002



2. Retrieve Knowledge (Vector Search)

0.84s | 0 chunks returned ⚠️



3. Answer Draft (GPT-4o)

2.18s | 1.2k in | 612 out | \$0.018



4. Tool: Get Account (Internal API)

0.73s | 200ms upstream



5. Policy Check (Claude Sonnet)

9.02s | 21k in | 850 out | \$0.142 ⚠️ Rate limit fallback



6. Compose Final (GPT-4o-mini)

0.91s | 410 in | 128 out | \$0.003



7. Log & Persist (Internal)

0.31s | DB write

With traces, you can finally answer:



Slow?

Step 5 (Claude fallback) took 9s due to rate limiting.



Expensive?

\$540/day spike is one user recursively calling the agent.



Wrong?

Retrieval returned 0 chunks, poisoning the downstream prompt.

Bottom line: Traces turn blind runs into explainable runs. You fix the right thing, faster.



LangSmith



Helicone



Arize



Langfuse

Why A wins:

Every problem in the scenario is the same root cause — you can't see the shape of a single agent run.

Slow? → Which of the 7 LLM calls was slow.

Expensive? → Which model on which step burned the tokens.

Wrong? → Which retrieval returned garbage that poisoned the next prompt.

None of those questions have answers without a trace. And LLM traces aren't HTTP traces — you need span-level attributes: model name, prompt tokens, completion tokens, cost, tool name, retrieval doc IDs, and the actual prompt/response payload. That's what LangSmith, Helicone, Arize, and Langfuse give you out of the box.

Once you have it: the 14-second span opens into 7 sub-spans. You see the one that took 9s was a Claude fallback triggered by a rate limit. The \$540/day spike is one user hammering an agent that recursively calls itself. The hallucinations correlate with retrievals that returned 0 chunks.

Answer 2

Why B is the senior engineer trap:

Evals answer a different question: *is quality trending up or down over time?* They're not a debugging tool. When prod is on fire, "78% correctness on last 200 samples" tells you nothing about the request that broke. Worse — evals without tracing are unfalsifiable. When the score drops, you have no way to drill into failing samples to see *why*. Evals are step 2. They ride on top of traces.

Answer 3

Why C treats the symptom:


"Spend is up 3x" is the same info you already had from the bill. Aggregate cost without span-level attribution just tells you the room is on fire without telling you which wall. You can't build per-agent, per-user cost attribution without instrumenting spans first. Budget alerts: yes, ship them. As your primary investment: you're solving the CFO's problem, not the on-call engineer's.

Answer 4

Why D isn't actually pragmatic:

"Own your data, no vendor lock-in" is the sentence that has cost more teams more months than almost any other. You'd need: schema for nested agent runs, trace tree visualizer, correlation IDs across async tool calls, sampling logic, retention, PII redaction, cost rollups. You just rebuilt Langfuse badly, in your spare time, while the incident is still open. Every serious LLM observability tool emits OpenTelemetry-compatible spans — switch vendors in a week if needed. Ship the vendor now, revisit self-hosting when you've earned that problem.

Designing Idempotency Keys for Payments



Your payment service charged a customer twice.

Same request. Different response code. Client retried. You had no dedup logic.

That's not a race condition. That's a missing design decision.

And the decision isn't whether to use idempotency keys. It's how you generate them.

Here's the scenario:

Your API processes payments. Clients can and will retry on network failures. Your downstream calls (Stripe, Kafka, wallet credit) cannot be undone once triggered. A 500 or timeout on your side looks identical to a "try again" from the client's perspective.

Which idempotency strategy do you go with?

A) Client-generated keys

Client sends a UUID with every request (Idempotency-Key:). Your server stores key → response in Redis with a 24h TTL. Duplicate request? Return cached response. No reprocessing.

B) Server-generated keys

Client calls POST /payments/init first. Server returns a one-time key. Client uses that key on the actual payment call. Server owns key lifecycle — rotation, expiry, scope.

C) Deterministic key from business fields

No UUID. Hash the fields that define uniqueness: sha256(customer_id + order_id + amount + currency). Same business intent = same key. No client coordination needed.

D) Outbox + idempotent consumer


Answer 1

✓ CORRECT ANSWER

A) Client-generated keys — ✓ GOLD STANDARD FOR EXTERNAL APIs

System Design Labs

This is how Stripe does it. How Braintree does it.

Reason: the client owns the intent, so the client owns the key.

A) Client-generated keys — ✓ GOLD STANDARD FOR EXTERNAL APIs

This is how Stripe does it. How Braintree does it. Reason: the client owns the intent, so the client owns the key.

The flow that actually works in production:

1. Write the key with status pending before processing (Redis SET NX)
2. Process the payment
3. Update key to completed with the full response body

Write the key first. Not after. If you crash between the payment succeeding and the key write — you've charged the card with no dedup record. Next retry charges again.

Key scoping matters: scope to (customer_id, key) not globally. TTL: 24h standard, 72h for aggressive retry clients. After expiry, same UUID = new request — document this.

Answer 2

B) Server-generated keys — RIGHT WHEN YOU DON'T TRUST CLIENTS

Two-step feels like overhead. It's actually a feature.

The init call gives you: pre-validation before money moves, audit trail at key issuance, protection against clients sending garbage UUIDs (v1 timestamps, sequential IDs, non-unique values), and full key lifecycle control server-side.

Where it shines: internal service-to-service APIs, regulated environments, anywhere key collision is a compliance issue.

Where it hurts: mobile clients on flaky networks now have two calls that can fail independently. Retry complexity goes up, not down.

Answer 3

C) Deterministic key — THE SENIOR ENGINEER TRAP

Looks clean. Falls apart on legitimate retries.

Same customer, same amount, same item — but a fresh order after a refund. Your hash matches the completed request. You return cached success for a brand new charge that should go through.

Second failure mode: field serialization drift across SDK versions. SDK v1 orders fields differently than v2. Different hash. Your dedup window disappears silently.

Use deterministic keys for internal job dedup — cron tasks, batch pipelines, where business intent is truly unique. Never for customer-facing payment flows.

Answer 4

D) Outbox + idempotent consumer —  STRONGEST FOR DISTRIBUTED SYSTEMS

The only option that makes the payment record and downstream event atomic.

The failure in A/B/C: Redis key write, DB write, and Stripe call are three separate operations. Crash between any two = inconsistent state.


The outbox fixes it:

1. Write payment record + outbox event in one DB transaction
2. Worker processes downstream (Stripe, Kafka) using the row ID as dedup token
3. Downstream consumers deduplicate on event ID

No separate key store. No Redis dependency. Exactly-once enforced at the DB layer.

The cost: outbox table + reliable worker + idempotent consumers. Real infrastructure. Overkill at 100 payments/day. The only defensible architecture at 10K TPS with regulatory requirements.

Architecting a SaaS Platform from Scratch



You're designing the architecture for a new SaaS platform from scratch.

It needs to handle 50,000 requests per second at launch. Multi-tenant. Real-time data. Global users. AI inference built in.

Your CTO gives you 3 months and a blank slate.

You've been thinking about this for a week. You finally sit down to draw the first box.

Here's what engineers actually debate in that room:

A) Start with a distributed microservices architecture — services for auth, billing, inference, and core domain from day one. Scale each independently.

B) Start with a modular monolith — one deployable unit with clean internal module boundaries. Extract services only when a specific boundary proves it needs it.

C) Start with serverless — Lambda/Cloud Functions for every endpoint, DynamoDB for state, SQS for async. Zero infra to manage, scales to zero between bursts.

D) Start with an event-driven architecture — Kafka as the backbone, all services communicate through events, no direct API calls between services from day one.

Answer 1


✓ CORRECT ANSWER



B) Modular Monolith — right

System Design Labs

The constraint isn't scale. It's the 3-month runway plus uncertainty about domain boundaries.



A modular monolith gives you:



Fast iteration — no network serialization overhead, no cross-service API contracts to version



Real module boundaries enforced in code (auth, billing, inference, core domain — no shared DB tables, no cross-module direct calls)




Single deploy, single observability surface, single transaction boundary



The option to extract services later — when you know exactly which boundary is real

Extraction Path Example: Billing



AI inference piece: inference latency is dominated by model I/O (50–500ms), not your application code. That's a reason to put it behind a clean internal interface — not a reason to make it a separate service on day one.



Modular now. Distributed later. Boundaries first, extraction second. Refactor, not rewrite.

B) Modular Monolith — right

The constraint isn't scale. It's the 3-month runway plus uncertainty about domain boundaries.

A modular monolith gives you:

- Fast iteration — no network serialization overhead, no cross-service API contracts to version
- Real module boundaries enforced in code (auth, billing, inference, core domain — no shared DB tables, no cross-module direct calls)
- Single deploy, single observability surface, single transaction boundary
- The option to extract services later — when you know exactly which boundary is real

When billing hits 10k RPS independently and needs separate scaling, you extract it. The boundary is already clean. The migration is a refactor, not a rewrite.

The AI inference piece specifically: inference latency is dominated by model I/O (50–500ms), not your application code. That's a reason to put it behind a clean internal interface — not a reason to make it a separate service on day one.

Answer 2

A) Microservices from day one — wrong

50,000 RPS sounds like you need microservices. It doesn't.

That's a solved load balancing problem, not a distribution problem.

What microservices from day one actually gives you:

- 6+ independently deployable services before you've found your real domain boundaries
- Distributed transactions across auth, billing, inference, and core domain
- Network hops on every request that used to be in-process calls
- 3 months of infra work instead of product

Amazon's "two-pizza teams" rule was for *existing* services with known boundaries. Not for new domains where you're still discovering what the boundaries are.

The boundary is the hard part. You don't know it yet.

Answer 3

C) Serverless — wrong (senior engineer trap)

"Scales to zero" sounds great until you price out 50k RPS sustained on Lambda.

The math:

- $50,000 \text{ req/s} \times 100\text{ms avg} = 5,000 \text{ concurrent Lambda executions}$
- At that scale you're paying more than a reserved EC2 cluster
- Cold starts are real and you promised "real-time data"
- AI inference in Lambda = cold start + memory limit + 15-min hard timeout = pain

Serverless wins for bursty, unpredictable event workloads. SaaS with consistent load and real-time requirements is not that workload.

Answer 4

D) Event-driven from day one — wrong

This one trips up engineers who've read about Kafka at Netflix and Uber.

Event-driven is the right answer — *for* Netflix and Uber, after they've already mapped their domain.

Starting with Kafka before you have a working system means:

- Every feature needs an event schema, a producer, and a consumer before it works
- Debugging "why didn't this event get processed" during week 2 of development
- Eventual consistency as the default when you don't yet know which operations need strong consistency
- Kafka operational overhead before you've validated the product

Event-driven is an optimization for a known system. Not a foundation for an unknown one.